

Data Structures

02

In This Edition

1	Overview	2
2	Master Complexity Table	2
3	The Data Structures	3
	3.1 Arrays & Dynamic Arrays	3
	3.2 Strings	4
	3.3 Hash Tables	5
	3.4 Linked Lists	7
	3.5 Stacks & Queues	9
	3.6 Trees	11
	3.7 Heaps / Priority Queues	13
	3.8 Tries (Prefix Trees)	15
	3.9 Graphs	17
	3.10 Union-Find / Disjoint Set	19
	3.11 LRU / LFU Cache Structures	21
	3.12 Segment Tree & Fenwick Tree (BIT)	23
	3.13 Bloom Filter	25
4	Choosing the Right Data Structure	26
5	Language Notes	26
	5.1 Python	26
	5.2 Java	27
	5.3 C++	28
6	Real-Life Analogies	28
7	Common Interview Questions	29
	7.1 Arrays	29
	7.2 Strings	29
	7.3 Hash Maps/Sets	29
	7.4 Linked Lists	30
	7.5 Stacks/Queues	30
	7.6 Trees	30
	7.7 Heaps	30
	7.8 Tries	30
	7.9 Graphs	30
	7.10 Union-Find	31
	7.11 LRU/LFU	31
	7.12 Segment Tree / BIT	31
8	Typical Mistakes Candidates Make	31
9	Revision Cheat Sheet	31
	9.1 Complexities to Memorize (Absolute Must)	32
	9.2 Decision Hooks (6 questions to ask every interview)	32
	9.3 Quick Pattern Mapping	32
	9.4 The 5 Structures You MUST Be Able to Code From Scratch	32
	9.5 Memory Aids	33
10	Visual Reference	34

Dense reference: mechanics, complexities, code, diagrams, and interview takeaways in one place.

1 Overview

Data structures are the substrate of every algorithm. Choosing the wrong one is the single most common reason candidates fail coding interviews — the algorithm is right, but $O(n)$ lookups in a list where a hash map would give $O(1)$ makes the solution too slow. This file is organized as a working reference: read top-to-bottom once, then use the complexity table and decision guide during revision.

Mental model for selection:

1. What operations do I need? (search, insert, delete, min/max, range, prefix?)
2. What are the frequency/order constraints? (random access? FIFO? LIFO? sorted?)
3. What are the acceptable complexities?
4. Then pick the structure.

2 Master Complexity Table

All complexities are average-case unless marked. W = worst case.

Data Structure	Access	Search	Insert	Delete	Space	Notes
Array (static)	$O(1)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	Insert/delete shifts elements
Dynamic Array (list)	$O(1)$	$O(n)$	$O(1)$ amort	$O(n)$	$O(n)$	Append amortized $O(1)$
String (immutable)	$O(1)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	Concat is $O(n)$; use StringBuilder
Hash Table	N/A	$O(1)$	$O(1)$	$O(1)$	$O(n)$	W: $O(n)$ with all collisions
Singly Linked List	$O(n)$	$O(n)$	$O(1)$ head	$O(n)$	$O(n)$	$O(1)$ if you have the node reference
Doubly Linked List	$O(n)$	$O(n)$	$O(1)$ ends	$O(1)^*$	$O(n)$	* $O(1)$ if node ref known
Stack	$O(n)$	$O(n)$	$O(1)$ push	$O(1)$ pop	$O(n)$	Top access $O(1)$
Queue	$O(n)$	$O(n)$	$O(1)$ enq	$O(1)$ deq	$O(n)$	Front access $O(1)$
Deque	$O(1)$ ends	$O(n)$	$O(1)$ ends	$O(1)$ ends	$O(n)$	Both ends $O(1)$
BST (balanced)	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(n)$	Unbalanced: $O(n)$ worst
BST (unbalanced)	$O(n)$ W	$O(n)$ W	$O(n)$ W	$O(n)$ W	$O(n)$	Degenerate = linked list
AVL Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(n)$	Strict balance; fast lookup
Red-Black Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(n)$	Looser balance; fast insert/delete
Binary Heap (min/max)	$O(1)$ top	$O(n)$	$O(\log n)$	$O(\log n)$	$O(n)$	Build heap: $O(n)$
Trie	$O(m)$	$O(m)$	$O(m)$	$O(m)$	$O(n*m)$	m = key length
Graph (adj list)	$O(V+E)$	$O(V+E)$	$O(1)$	$O(E)$	$O(V+E)$	Sparse graphs
Graph (adj matrix)	$O(1)$	$O(1)$ edge	$O(1)$	$O(1)$	$O(V^2)$	Dense graphs
Union-Find	$O(\alpha(n))$	$O(\alpha(n))$	$O(\alpha(n))$	-	$O(n)$	α = inverse Ackermann, $\approx O(1)$

Data Structure	Access	Search	Insert	Delete	Space	Notes
Segment Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(n)$	Range queries + point updates
Fenwick Tree (BIT)	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(n)$	Simpler than segment tree
LRU Cache	$O(1)$	$O(1)$	$O(1)$	$O(1)$	$O(n)$	Hash map + doubly linked list
Bloom Filter	N/A	$O(k)$	$O(k)$	N/A	$O(m)$	k =hash funcs, probabilistic

3 The Data Structures

3.1 Arrays & Dynamic Arrays

What it is: Contiguous block of memory where each element sits at a fixed offset from the base address. Random access is $O(1)$ because $\text{address} = \text{base} + \text{index} * \text{element_size}$.

Internal mechanics:

```
Static Array (size=6):
Index:  [0] [1] [2] [3] [4] [5]
Value:  [10] [20] [30] [40] [50] [60]
        ^               ^
        base addr       base + 5*sizeof(int)

Dynamic Array (Python list) – amortized resize:
size=4, capacity=4:  [1][2][3][4][ ][ ][ ][ ] ← capacity doubles on overflow
Append 5 → capacity=8:[1][2][3][4][5][ ][ ][ ]
Append cost analysis (n appends):
- n-1 cheap appends:  $O(1)$  each
-  $\sim \log(n)$  doublings: cost  $1+2+4+\dots+n = 2n$  copies total
- Amortized per append:  $2n / n = O(1)$  ✓
```

Complexity:

Op	Best	Avg	Worst
Access	$O(1)$	$O(1)$	$O(1)$
Search	$O(1)$	$O(n)$	$O(n)$
Append	$O(1)$	$O(1)$	$O(n)$
Insert	$O(1)$	$O(n)$	$O(n)$
Delete	$O(1)$	$O(n)$	$O(n)$

When to use:

- › Random access by index is needed
- › Data is mostly read; writes are appends
- › Cache locality matters (iteration speed)

When NOT to use:

- › Frequent mid-list insertions/deletions (use linked list or deque)
- › Need fast search (use hash map or BST)

Key Python patterns:

```

# Two-pointer on sorted array
left, right = 0, len(arr) - 1
while left < right:
    s = arr[left] + arr[right]
    if s == target: return [left, right]
    elif s < target: left += 1
    else: right -= 1

# Sliding window
window_sum = sum(arr[:k])
max_sum = window_sum
for i in range(k, len(arr)):
    window_sum += arr[i] - arr[i - k]
    max_sum = max(max_sum, window_sum)

# In-place reversal
arr[i], arr[j] = arr[j], arr[i]

```

Interview takeaway: Arrays are the default. When the problem says "return sorted" or "in-place", think array + two pointers. When you need to insert/delete mid-sequence often, question whether an array is right.

Real interview use-case: Two Sum (hash map), Sliding Window Maximum, Kadane's (max subarray), rotate array in-place, merge sorted arrays.

3.2 Strings

What it is: Sequence of characters. In Python and Java, strings are **immutable** — every concatenation creates a new object. In C++, `std::string` is mutable but each `+=` may reallocate.

Internal mechanics:

```

Python string "hello":
h | e | l | l | o
Each character is a Unicode code point (Python 3: variable width internally)
id("hello" + "x") != id("hello") ← new object every concat

String concatenation cost in a loop (naïve):
for i in range(n): result += s[i]
Iteration 1: copy 1 char → O(1)
Iteration 2: copy 2 chars → O(2)
...
Iteration n: copy n chars → O(n)
Total: O(1+2+...+n) = O(n²) ← DISASTER for large n

String Builder pattern:
parts = []
for c in chars: parts.append(c) ← O(1) amortized each
result = "".join(parts) ← O(n) one pass
Total: O(n) ✓

```

Complexity:

Op	Complexity	Note
Index access	$O(1)$	
Slice <code>s[i:j]</code>	$O(j-i)$	Creates new string
Concatenate	$O(n+m)$	Both strings copied
Find/search	$O(n*m)$	Naive; KMP/Z gives $O(n+m)$
Comparison	$O(n)$	Character by character

Join

O(n)

One pass

Key Python patterns:

```

# Always join, never += in loop
result = "".join(list_of_chars)

# Reverse a string
rev = s[::-1]

# Check palindrome
def is_palindrome(s):
    s = s.lower()
    l, r = 0, len(s) - 1
    while l < r:
        while l < r and not s[l].isalnum(): l += 1
        while l < r and not s[r].isalnum(): r -= 1
        if s[l] != s[r]: return False
        l += 1; r -= 1
    return True

# Anagram check
from collections import Counter
Counter(s) == Counter(t)

# Sliding window on string (distinct chars)
from collections import defaultdict
freq = defaultdict(int)
l = 0
for r, c in enumerate(s):
    freq[c] += 1
    while len(freq) > k:
        freq[s[l]] -= 1
        if freq[s[l]] == 0: del freq[s[l]]
        l += 1

```

Interview takeaway: Never concatenate strings in a loop. Use `"".join(parts)`. For pattern matching, think sliding window + frequency map. Anagram/permutation problems → character frequency counter.

Real interview use-case: Longest substring without repeating characters, group anagrams, valid palindrome, minimum window substring, encode/decode strings.

3.3 Hash Tables

What it is: Maps keys to values in O(1) average by computing `hash(key) % capacity` to find a bucket, then storing the key-value pair there.

Internal mechanics:

Hash Table with Chaining (load factor = 0.75):

```

Buckets:
[0] → None
[1] → ("apple", 5) → ("grape", 2) → None ← collision! same bucket
[2] → ("banana", 3) → None
[3] → None
[4] → ("cherry", 7) → None
...

hash("apple") % 8 = 1
hash("grape") % 8 = 1 ← collision → chain extends

```

```
hash("banana") % 8 = 2
```

Open Addressing (linear probing):

```
[0]: empty
[1]: ("apple", 5)
[2]: ("grape", 2) ← collision at 1, probe to 2
[3]: ("banana", 3)
...
```

Load Factor (α) = $n / \text{capacity}$

$\alpha < 0.7$: few collisions, $O(1)$ expected

$\alpha > 0.9$: many collisions $\rightarrow O(n)$ worst

$\rightarrow$ resize (rehash all) when α threshold exceeded

Python dict resizes at 2/3 full $\rightarrow$ rehash to 2x capacity

Good Hash Function Properties:

1. Deterministic (same key $\rightarrow$ same hash)
2. Uniform distribution (minimize collisions)
3. Fast to compute
4. Avalanche effect (small key change $\rightarrow$ big hash change)

Collision resolution comparison:

Strategy	Pro	Con
Chaining	Simple, works at high load	Extra memory for pointers
Linear probing	Cache-friendly	Clustering degrades performance
Quadratic probe	Less clustering	May not probe all slots
Double hashing	Best distribution	Two hash functions needed

Complexity:

Op	Average	Worst
Search	$O(1)$	$O(n)$
Insert	$O(1)$	$O(n)$
Delete	$O(1)$	$O(n)$

When to use:

- › $O(1)$ lookup by key
- › Counting frequencies
- › Caching / memoization
- › Deduplication

When NOT to use:

- › Need sorted order (use BST/sorted array)
- › Keys not hashable (use tree map)
- › Memory is very tight

Key Python patterns:

```
from collections import defaultdict, Counter

# Frequency count
freq = Counter("anagram")          # Counter({'a': 3, 'n': 1, ...})
```

```
# Default dict avoids KeyError
graph = defaultdict(list)
graph['A'].append('B')

# Two Sum O(n)
def two_sum(nums, target):
    seen = {}
    for i, n in enumerate(nums):
        complement = target - n
        if complement in seen:
            return [seen[complement], i]
        seen[n] = i

# Grouping anagrams
from collections import defaultdict
def group_anagrams(strs):
    groups = defaultdict(list)
    for s in strs:
        groups[tuple(sorted(s))].append(s)
    return list(groups.values())
```

Interview takeaway: Hash map is your first tool when you need $O(1)$ lookup. "Two pointer $O(n^2)$? Can we trade space for time?" → hash map. Watch out: dict keys must be hashable (no lists as keys; use tuples).

Real interview use-case: Two Sum, subarray sum equals k, longest consecutive sequence, top K frequent elements, valid sudoku.

3.4 Linked Lists

What it is: Sequence of nodes where each node holds a value and a pointer to the next (singly) or both next and previous (doubly) nodes. Non-contiguous memory.

Internal mechanics:

```
Singly Linked List:
head
↓
[1|•]→[2|•]→[3|•]→[4|•]→None
node  next pointer

Doubly Linked List:
head                                tail
↓                                ↓
None←[•|1|•]↔[•|2|•]↔[•|3|•]↔[•|4|•]→None
prev val next

Insert at head (O(1)):
new_node.next = head
head = new_node

Insert at tail (O(1) with tail pointer):
tail.next = new_node
tail = new_node

Delete node (given reference, doubly linked, O(1)):
node.prev.next = node.next
node.next.prev = node.prev

Delete node (singly linked, given reference only, O(n)):
Must traverse from head to find predecessor
→ trick: copy next node's value into current, delete next
```

Complexity:

Op	Singly LL	Doubly LL
Access by index	$O(n)$	$O(n)$
Search	$O(n)$	$O(n)$
Insert at head	$O(1)$	$O(1)$
Insert at tail	$O(n)/O(1)$	$O(1)$
Insert at middle	$O(n)$	$O(n)$
Delete at head	$O(1)$	$O(1)$
Delete by ref	$O(n)$	$O(1)$

When to use:

- › Frequent insertions/deletions at head or known positions
- › Implementing stacks, queues, LRU cache
- › When you don't need random access

When NOT to use:

- › Random access needed (use array)
- › Memory is tight (each node has pointer overhead)
- › Cache performance critical (non-contiguous = cache misses)

Key Python patterns:

```

class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

# Reverse a linked list (iterative)
def reverse(head):
    prev = None
    curr = head
    while curr:
        nxt = curr.next
        curr.next = prev
        prev = curr
        curr = nxt
    return prev

# Floyd's cycle detection
def has_cycle(head):
    slow = fast = head
    while fast and fast.next:
        slow = slow.next
        fast = fast.next.next
    if slow is fast:
        return True
    return False

# Find middle (slow/fast pointers)
def find_middle(head):
    slow = fast = head
    while fast and fast.next:
        slow = slow.next
        fast = fast.next.next
    return slow

```

```
# Dummy node pattern (simplifies edge cases)
def delete_nth_from_end(head, n):
    dummy = ListNode(0, head)
    fast = slow = dummy
    for _ in range(n + 1):
        fast = fast.next
    while fast:
        fast = fast.next
        slow = slow.next
    slow.next = slow.next.next
    return dummy.next
```

Interview takeaway: Always use a dummy head node to avoid null checks at head. Slow/fast pointer solves cycle detection, middle finding, and nth-from-end in one pass. Linked list problems are usually just pointer manipulation — draw it out.

Real interview use-case: Reverse linked list, merge two sorted lists, detect cycle, reorder list, LRU cache.

3.5 Stacks & Queues

What it is:

- › **Stack:** LIFO (Last In, First Out). Like a stack of plates.
- › **Queue:** FIFO (First In, First Out). Like a line at a store.
- › **Deque:** Double-ended queue; insert/delete at both ends $O(1)$.
- › **Monotonic stack/queue:** Stack/queue that maintains monotone order for range queries.
- › **Circular buffer:** Fixed-size ring buffer for queue; avoids shifting.

Internal mechanics:

```
Stack (LIFO):
Push 1,2,3: [1]→[2]→[3] ← top is 3
Pop:        [1]→[2]      ← returns 3

Queue (FIFO) using circular buffer:
Capacity=5, head=0, tail=0
Enqueue 1,2,3:
[1][2][3][ ][ ] head=0, tail=3
Deque:
[ ][2][3][ ][ ] head=1, tail=3
Enqueue 4,5:
[ ][2][3][4][5] head=1, tail=0 (wrapped!)
→ indices mod capacity: (tail+1) % cap

Monotonic Stack (next greater element):
arr = [2, 1, 5, 3, 4]
Stack tracks indices of elements waiting for their "next greater"
i=0: push 0      stack=[0]      (arr[0]=2)
i=1: 1<2, push 1  stack=[0,1]    (arr[1]=1)
i=2: 5>1, pop 1→ans[1]=5, 5>2 pop 0→ans[0]=5, push 2
      stack=[2]
i=3: 3<5, push 3  stack=[2,3]
i=4: 4>3, pop 3→ans[3]=4, 4<5 push 4
      stack=[2,4]
Remaining: no greater element → -1
ans = [5, 5, -1, 4, -1]

Deque (collections.deque in Python):
appendleft / popleft →  $O(1)$ 
append / pop        →  $O(1)$ 
```

Complexity:

Op	Stack	Queue	Deque	Monotonic Stack
Push/Enqueue	O(1)	O(1)	O(1)	O(1) amort
Pop/Dequeue	O(1)	O(1)	O(1)	O(1) amort
Peek (top/front)	O(1)	O(1)	O(1)	O(1)
Search	O(n)	O(n)	O(n)	O(n)

When to use:

- › Stack: DFS, expression parsing, undo operations, backtracking
- › Queue: BFS, task scheduling, sliding window with deque
- › Monotonic stack: next greater/smaller element, largest rectangle in histogram
- › Deque: sliding window maximum, palindrome check

Key Python patterns:

```

from collections import deque

# Stack (use list, pop from end)
stack = []
stack.append(x)    # push
stack.pop()        # pop O(1)
stack[-1]          # peek

# Queue (use deque for O(1) dequeue)
q = deque()
q.append(x)         # enqueue
q.popleft()         # dequeue O(1)

# BFS template
def bfs(root):
    q = deque([root])
    while q:
        node = q.popleft()
        for child in node.children:
            q.append(child)

# Monotonic stack – next greater element
def next_greater(nums):
    n = len(nums)
    ans = [-1] * n
    stack = [] # indices
    for i, num in enumerate(nums):
        while stack and nums[stack[-1]] < num:
            idx = stack.pop()
            ans[idx] = num
        stack.append(i)
    return ans

# Sliding window maximum (monotonic deque)
def sliding_max(nums, k):
    dq = deque() # indices, decreasing values
    res = []
    for i, num in enumerate(nums):
        while dq and nums[dq[-1]] < num:
            dq.pop()
        dq.append(i)
        if dq[0] == i - k: # out of window
            dq.popleft()

```

```

    if i >= k - 1:
        res.append(nums[dq[0]])
    return res

```

Interview takeaway: Monotonic stack is the pattern for "next greater/smaller element" and "largest rectangle" problems. BFS always uses a queue. For sliding window max/min, use a monotonic deque. Never use `list.pop(0)` for queues — it's $O(n)$; use `deque.popleft()`.

Real interview use-case: Valid parentheses, min stack, daily temperatures, largest rectangle in histogram, sliding window maximum.

3.6 Trees

What it is: Hierarchical, acyclic connected graph. Root at top, leaves at bottom.

Internal mechanics:

Binary Tree (each node: up to 2 children):



BST (Binary Search Tree): $\text{left} < \text{root} < \text{right}$



BST invariant: for any node N:

- All nodes in left subtree $< N.val$
- All nodes in right subtree $> N.val$
- In-order traversal gives sorted sequence!

BST Search:

Find 6: start at 8 → $6 < 8$ go left → 3 → $6 > 3$ go right → 6 ✓
 $O(h)$ where h = height. Balanced: $h = \log n$. Skewed: $h = n$.

Balanced Tree (AVL) – height constraint:

Balance factor = $\text{height}(\text{left}) - \text{height}(\text{right})$

Must be in $\{-1, 0, 1\}$ for every node

Rotation on violation:

- Left-Left: right rotate
- Right-Right: left rotate
- Left-Right: left rotate then right rotate
- Right-Left: right rotate then left rotate

Red-Black Tree (less strict, used in Java TreeMap):

Properties:

1. Every node is red or black
 2. Root is black
 3. Red nodes have black children
 4. All paths from node to null have same # black nodes
- Height $\leq 2 \cdot \log(n+1)$ guaranteed

Tree Traversals:

Pre-order:	root → left → right	[1,2,4,5,3,6]
In-order:	left → root → right	[4,2,5,1,3,6] ← sorted for BST!
Post-order:	left → right → root	[4,5,2,6,3,1]
Level-order:	BFS level by level	[1,2,3,4,5,6]

Complexity:

Op	Balanced BST	Unbalanced BST	AVL / RB Tree
Search	$O(\log n)$	$O(n)$	$O(\log n)$
Insert	$O(\log n)$	$O(n)$	$O(\log n)$
Delete	$O(\log n)$	$O(n)$	$O(\log n)$
Min/Max	$O(\log n)$	$O(n)$	$O(\log n)$
In-order	$O(n)$	$O(n)$	$O(n)$

When to use:

- › BST: sorted data with dynamic insertions/deletions
- › AVL/RB: guaranteed $O(\log n)$ when data can be adversarial
- › Any tree: hierarchical relationships, recursive decomposition

Key Python patterns:

```

class TreeNode:
    def __init__(self, val=0, left=None, right=None):
        self.val = val
        self.left = left
        self.right = right

# In-order traversal (recursive)
def inorder(root):
    if not root: return []
    return inorder(root.left) + [root.val] + inorder(root.right)

# In-order traversal (iterative – important for interviews)
def inorder_iter(root):
    result, stack = [], []
    curr = root
    while curr or stack:
        while curr:
            stack.append(curr)
            curr = curr.left
        curr = stack.pop()
        result.append(curr.val)
        curr = curr.right
    return result

# Height of tree
def height(root):
    if not root: return 0
    return 1 + max(height(root.left), height(root.right))

# Check if BST is valid
def is_valid_bst(root, lo=float('-inf'), hi=float('inf')):
    if not root: return True
    if not (lo < root.val < hi): return False
    return (is_valid_bst(root.left, lo, root.val) and
            is_valid_bst(root.right, root.val, hi))

# LCA (Lowest Common Ancestor) in BST
def lca_bst(root, p, q):
    if p.val < root.val and q.val < root.val:
        return lca_bst(root.left, p, q)
    if p.val > root.val and q.val > root.val:
        return lca_bst(root.right, p, q)
    return root

```

```
# Level order traversal
from collections import deque
def level_order(root):
    if not root: return []
    result, q = [], deque([root])
    while q:
        level = []
        for _ in range(len(q)):
            node = q.popleft()
            level.append(node.val)
            if node.left: q.append(node.left)
            if node.right: q.append(node.right)
        result.append(level)
    return result
```

Interview takeaway: BST in-order = sorted. Use recursion with min/max bounds for BST validation, not just $\text{left} < \text{root} < \text{right}$ (common bug). Diameter of tree = longest path through any node = $\max(\text{left_depth} + \text{right_depth})$ over all nodes.

Real interview use-case: Validate BST, lowest common ancestor, diameter of binary tree, serialize/deserialize tree, path sum, kth smallest in BST.

3.7 Heaps / Priority Queues

What it is: A complete binary tree stored as an array, where each parent satisfies the heap property (min-heap: $\text{parent} \leq \text{children}$; max-heap: $\text{parent} \geq \text{children}$). Top element is always the min (or max).

Internal mechanics:

Min-Heap (array representation):

```

      1
     /\
    3  5
   /\ /\
  7 4 8 9
```

Array: [1, 3, 5, 7, 4, 8, 9]

Index: 0 1 2 3 4 5 6

Index relationships:

Parent(i) = (i-1) // 2

Left child(i) = 2*i + 1

Right child(i) = 2*i + 2

Sift Up (after insert at end):

Insert 2: append → [1,3,5,7,4,8,9,2]

Compare 2 with parent (index 3): $\text{arr}[3]=7 > 2 \rightarrow \text{swap}$

→ [1,3,5,2,4,8,9,7]

Compare 2 with parent (index 1): $\text{arr}[1]=3 > 2 \rightarrow \text{swap}$

→ [1,2,5,3,4,8,9,7]

Compare 2 with parent (index 0): $\text{arr}[0]=1 < 2 \rightarrow \text{stop} \checkmark$

Sift Down (after extracting root – put last element at root):

Extract 1: put 7 at root → [7,2,5,3,4,8,9]

Compare 7 with children 2,5: min child=2 → swap

→ [2,7,5,3,4,8,9]

Compare 7 with children 3,4: min child=3 → swap

→ [2,3,5,7,4,8,9]

7 has no children → stop $\checkmark$

Build Heap (heapify) from arbitrary array – $O(n)$:

Start sifting down from last non-leaf ($n/2 - 1$) to root

Counterintuitive: $O(n)$ not $O(n \log n)$ because most nodes are near leaves

Complexity:

Op	Time
Get min/max	$O(1)$
Insert	$O(\log n)$
Extract min	$O(\log n)$
Build heap	$O(n)$
Search	$O(n)$
Heapsort	$O(n \log n)$

When to use:

- › Get the min/max repeatedly (priority queues)
- › Top K elements (use heap of size K)
- › Merge K sorted lists
- › Dijkstra's shortest path
- › Median maintenance (two heaps)

Key Python patterns:

```
import heapq

# Min-heap (default)
heap = []
heapq.heappush(heap, 3)
heapq.heappush(heap, 1)
heapq.heappush(heap, 2)
heapq.heappop(heap) # returns 1

# Max-heap: negate values
max_heap = []
heapq.heappush(max_heap, -3)
heapq.heappush(max_heap, -1)
-heapq.heappop(max_heap) # returns 3

# Build heap in O(n)
nums = [3, 1, 4, 1, 5, 9]
heapq.heapify(nums)

# Top K elements
def top_k(nums, k):
    return heapq.nlargest(k, nums) # O(n log k)

# Top K using min-heap of size k (space efficient)
def top_k_heap(nums, k):
    heap = []
    for n in nums:
        heapq.heappush(heap, n)
        if len(heap) > k:
            heapq.heappop(heap)
    return sorted(heap, reverse=True)

# Merge K sorted lists
def merge_k_lists(lists):
    heap = []
    for i, lst in enumerate(lists):
        if lst:
            heapq.heappush(heap, (lst[0], i, 0))
    result = []
```

```

while heap:
    val, i, j = heapq.heappop(heap)
    result.append(val)
    if j + 1 < len(lists[i]):
        heapq.heappush(heap, (lists[i][j+1], i, j+1))
return result

# Median maintenance (two heaps)
class MedianFinder:
    def __init__(self):
        self.lo = [] # max-heap (negate)
        self.hi = [] # min-heap
    def add_num(self, num):
        heapq.heappush(self.lo, -num)
        heapq.heappush(self.hi, -heapq.heappop(self.lo))
        if len(self.hi) > len(self.lo):
            heapq.heappush(self.lo, -heapq.heappop(self.hi))
    def find_median(self):
        if len(self.lo) > len(self.hi): return -self.lo[0]
        return (-self.lo[0] + self.hi[0]) / 2.0

```

Interview takeaway: Python only has min-heap via `heapq`; negate values for max-heap. `heapify` is $O(n)$. Top K smallest: use max-heap of size K and push out the max when size exceeds K. Top K largest: use min-heap of size K.

Real interview use-case: Kth largest element, merge K sorted lists, task scheduler, find median from data stream, Dijkstra's.

3.8 Tries (Prefix Trees)

What it is: A tree where each path from root to leaf spells a word, and each node represents one character. Shared prefixes share nodes.

Internal mechanics:

Trie storing: ["cat", "car", "card", "care", "bat"]



(*) = is_end marker

Node structure:

```

children: dict[char → TrieNode] (or array[26])
is_end: bool

```

Insert "card":

```

root → c (create if absent)
      → a (create if absent)
      → r (exists, reuse)
      → d (create)
mark d.is_end = True

```

Search "care":

```

root → c → a → r → e → is_end? True ✓

```


StartsWith "ca":
 root → c → a → exists? True ✓ (prefix search)

Space: each node uses $O(\text{ALPHA})$ where $\text{ALPHA}=26$ for lowercase
 n strings of avg length m: $O(n*m)$ nodes worst case
 (better than storing all strings if many shared prefixes)

Complexity (m = length of key):

Op	Time	Space
Insert	$O(m)$	$O(m)$
Search	$O(m)$	$O(1)$
StartsWith	$O(m)$	$O(1)$
Delete	$O(m)$	$O(1)$

When to use:

- › Autocomplete / prefix search
- › Spell checking
- › IP routing tables
- › Word search in a grid

When NOT to use:

- › Simple key-value lookup (use hash map)
- › Keys are not strings or have no shared prefixes

Key Python patterns:

```

class TrieNode:
    def __init__(self):
        self.children = {}
        self.is_end = False

class Trie:
    def __init__(self):
        self.root = TrieNode()

    def insert(self, word: str) -> None:
        node = self.root
        for ch in word:
            if ch not in node.children:
                node.children[ch] = TrieNode()
            node = node.children[ch]
        node.is_end = True

    def search(self, word: str) -> bool:
        node = self.root
        for ch in word:
            if ch not in node.children:
                return False
            node = node.children[ch]
        return node.is_end

    def starts_with(self, prefix: str) -> bool:
        node = self.root
        for ch in prefix:

```

```

        if ch not in node.children:
            return False
        node = node.children[ch]
    return True

# Word search II (trie + DFS backtracking)
# Build trie from word list, then DFS on board
# When we reach a trie end node, we found a word

```

Interview takeaway: Trie is the answer to "find all words with prefix X" or "word search on a grid with a word list". For word search II (board + word list), build a trie from words and run DFS on the board — this prunes early when no prefix matches.

Real interview use-case: Implement trie, word search II, replace words, design search autocomplete, map sum pairs.

3.9 Graphs

What it is: A set of vertices (nodes) connected by edges. Can be directed/undirected, weighted/unweighted, cyclic/acyclic.

Internal mechanics:

```

Graph: 4 vertices, 4 edges (undirected)
0 --- 1
|   / |
|  /  |
| /   |
2 --- 3

Adjacency List (preferred for sparse graphs):
{
  0: [1, 2],
  1: [0, 2, 3],
  2: [0, 1, 3],
  3: [1, 2]
}
Space: O(V + E)
Check edge (u,v): O(degree(u))
Get neighbors: O(degree(u))

Adjacency Matrix (preferred for dense graphs):
  0  1  2  3
0 [ 0  1  1  0 ]
1 [ 1  0  1  1 ]
2 [ 1  1  0  1 ]
3 [ 0  1  1  0 ]
Space: O(V2)
Check edge (u,v): O(1)
Get neighbors: O(V)

Weighted adjacency list:
{0: [(1, 4), (2, 1)], ...} ← (neighbor, weight)

Types:
Undirected: edges go both ways
Directed (Digraph): edges have direction
DAG: Directed Acyclic Graph (used in topological sort)
Tree: connected undirected graph with V-1 edges

```

Complexity:

Op	Adj List	Adj Matrix
Space	$O(V+E)$	$O(V^2)$

Op	Adj List	Adj Matrix
Add vertex	$O(1)$	$O(V^2)$
Add edge	$O(1)$	$O(1)$
Remove edge	$O(E)$	$O(1)$
Check edge (u,v)	$O(\text{degree})$	$O(1)$
Get all neighbors	$O(\text{degree})$	$O(V)$
BFS/DFS	$O(V+E)$	$O(V^2)$
Dijkstra	$O((V+E)\log V)$	$O(V^2\log V)$

When to use:

- › Adj List: default, sparse graphs, BFS/DFS heavy
- › Adj Matrix: dense graphs, need $O(1)$ edge existence check, Floyd-Warshall

Key Python patterns:

```

from collections import defaultdict, deque

# Build graph
graph = defaultdict(list)
for u, v in edges:
    graph[u].append(v)
    graph[v].append(u) # undirected

# BFS (shortest path in unweighted graph)
def bfs(graph, start, end):
    visited = {start}
    q = deque([(start, 0)])
    while q:
        node, dist = q.popleft()
        if node == end: return dist
        for nei in graph[node]:
            if nei not in visited:
                visited.add(nei)
                q.append((nei, dist + 1))
    return -1

# DFS (iterative)
def dfs(graph, start):
    visited = set()
    stack = [start]
    while stack:
        node = stack.pop()
        if node in visited: continue
        visited.add(node)
        for nei in graph[node]:
            stack.append(nei)

# DFS (recursive)
def dfs_rec(node, visited, graph):
    visited.add(node)
    for nei in graph[node]:
        if nei not in visited:
            dfs_rec(nei, visited, graph)

# Topological sort (Kahn's algorithm, BFS-based)
def topo_sort(n, edges):

```

```

graph = defaultdict(list)
in_degree = [0] * n
for u, v in edges:
    graph[u].append(v)
    in_degree[v] += 1
q = deque(i for i in range(n) if in_degree[i] == 0)
order = []
while q:
    node = q.popleft()
    order.append(node)
    for nei in graph[node]:
        in_degree[nei] -= 1
        if in_degree[nei] == 0:
            q.append(nei)
return order if len(order) == n else [] # empty = cycle exists

# Dijkstra's (single source shortest path)
import heapq
def dijkstra(graph, start):
    dist = {node: float('inf') for node in graph}
    dist[start] = 0
    heap = [(0, start)]
    while heap:
        d, node = heapq.heappop(heap)
        if d > dist[node]: continue
        for nei, weight in graph[node]:
            new_d = d + weight
            if new_d < dist[nei]:
                dist[nei] = new_d
                heapq.heappush(heap, (new_d, nei))
    return dist

```

Interview takeaway: BFS for shortest path in unweighted graph. Dijkstra for weighted (non-negative). Detect cycle in directed graph: DFS with in-stack tracking. Topological sort only valid for DAGs. "Number of islands" = number of DFS/BFS calls on unvisited cells.

Real interview use-case: Number of islands, clone graph, course schedule, network delay time, Pacific Atlantic water flow, word ladder.

3.10 Union-Find / Disjoint Set

What it is: Data structure to track which elements belong to the same connected component. Supports two operations: find (which component?) and union (merge components).

Internal mechanics:

```

Initial state (6 elements, each its own component):
parent = [0, 1, 2, 3, 4, 5]
rank   = [0, 0, 0, 0, 0, 0]

Union(0, 1): find roots → 0,1; rank equal → make 1 child of 0
parent = [0, 0, 2, 3, 4, 5]
rank   = [1, 0, 0, 0, 0, 0]

Union(2, 3):
parent = [0, 0, 2, 2, 4, 5]

Union(0, 2): find roots → 0, 2; rank[0]=1 > rank[2]=0 → 2 under 0
parent = [0, 0, 0, 2, 4, 5] ← 2's parent is now 0

Path Compression during find(3):
find(3): parent[3]=2, parent[2]=0, parent[0]=0 → root=0
Compress: set parent[3]=0, parent[2]=0 directly

```

```
parent = [0, 0, 0, 0, 4, 5] ← next find(3) is 0(1)
```

After path compression and union by rank:
 find and union are effectively $O(\alpha(n)) \approx O(1)$
 where α = inverse Ackermann function
 $\alpha(n) < 5$ for any practical n

Visual after several unions:

Component 1: {0,1,2,3} ← tree rooted at 0

Component 2: {4}

Component 3: {5}

```

  0
 /|\
1 2 3 ← after path compression, flat tree

```

Complexity:

Op	Naive	Union by Rank	+ Path Compression
Find	$O(n)$	$O(\log n)$	$O(\alpha(n)) \approx O(1)$
Union	$O(n)$	$O(\log n)$	$O(\alpha(n)) \approx O(1)$
Space	$O(n)$	$O(n)$	$O(n)$

When to use:

- › Connected components (dynamic)
- › Detect cycles in undirected graph
- › Kruskal's minimum spanning tree
- › Grouping problems ("same group as")

Key Python patterns:

```

class UnionFind:
    def __init__(self, n):
        self.parent = list(range(n))
        self.rank = [0] * n
        self.components = n

    def find(self, x):
        if self.parent[x] != x:
            self.parent[x] = self.find(self.parent[x]) # path compression
        return self.parent[x]

    def union(self, x, y):
        rx, ry = self.find(x), self.find(y)
        if rx == ry: return False # already connected
        # Union by rank
        if self.rank[rx] < self.rank[ry]:
            rx, ry = ry, rx
        self.parent[ry] = rx
        if self.rank[rx] == self.rank[ry]:
            self.rank[rx] += 1
        self.components -= 1
        return True

    def connected(self, x, y):
        return self.find(x) == self.find(y)

# Example: count components in graph
def count_components(n, edges):

```

```

uf = UnionFind(n)
for u, v in edges:
    uf.union(u, v)
return uf.components

# Detect cycle in undirected graph
def has_cycle(n, edges):
    uf = UnionFind(n)
    for u, v in edges:
        if not uf.union(u, v): # already in same component
            return True
    return False

# Kruskal's MST
def kruskal(n, edges):
    uf = UnionFind(n)
    edges.sort(key=lambda e: e[2]) # sort by weight
    mst_cost = 0
    for u, v, w in edges:
        if uf.union(u, v):
            mst_cost += w
    return mst_cost

```

Interview takeaway: Always implement both path compression AND union by rank. Without both, you don't get the near- $O(1)$ guarantee. Union-Find cannot efficiently handle edge deletion — if you need that, use Link-Cut Trees (not interview scope).

Real interview use-case: Number of provinces, redundant connection, accounts merge, satisfiability of equality equations.

3.11 LRU / LFU Cache Structures

What it is:

- › **LRU (Least Recently Used):** Evicts the item not accessed for the longest time. $O(1)$ get and put.
- › **LFU (Least Frequently Used):** Evicts the item used fewest times (break ties by LRU). $O(1)$ all operations.

Internal mechanics:

```

LRU Cache (capacity=3):
Internals: HashMap + Doubly Linked List
Map: key → node ( $O(1)$  access)
DLL: most recent at front, least recent at tail

get(1): move node 1 to front
put(4): if full, remove tail (LRU), insert at front, update map

State after put(1), put(2), put(3):
[3] ↔ [2] ↔ [1] ← tail (LRU)
  ↑
  head (MRU)

get(1): move 1 to head
[1] ↔ [3] ↔ [2] ← tail (LRU)

put(4): evict tail (2), add 4 at head
[4] ↔ [1] ↔ [3]

LFU Cache (capacity=3):
Internals: HashMap(key→(val,freq)) + HashMap(freq→DLL of keys)
          + min_freq tracker
On access: increment freq, move to freq+1 list
On evict: remove LRU from min_freq list
Trickier: when min_freq list empties, min_freq++

```

Complexity:

Op	LRU	LFU
get	O(1)	O(1)
put	O(1)	O(1)

Key Python patterns:

```

from collections import OrderedDict

# LRU using OrderedDict (move_to_end trick)
class LRUCache:
    def __init__(self, capacity: int):
        self.cap = capacity
        self.cache = OrderedDict() # key → value, ordered by access time

    def get(self, key: int) -> int:
        if key not in self.cache: return -1
        self.cache.move_to_end(key) # mark as most recently used
        return self.cache[key]

    def put(self, key: int, value: int) -> None:
        if key in self.cache:
            self.cache.move_to_end(key)
        self.cache[key] = value
        if len(self.cache) > self.cap:
            self.cache.popitem(last=False) # remove LRU (first item)

# LRU from scratch (HashMap + DLL) — what interviewers actually want
class DLLNode:
    def __init__(self, key=0, val=0):
        self.key = key
        self.val = val
        self.prev = self.next = None

class LRUCacheManual:
    def __init__(self, capacity):
        self.cap = capacity
        self.cache = {}
        self.head = DLLNode() # dummy head (MRU side)
        self.tail = DLLNode() # dummy tail (LRU side)
        self.head.next = self.tail
        self.tail.prev = self.head

    def _remove(self, node):
        node.prev.next = node.next
        node.next.prev = node.prev

    def _insert_front(self, node):
        node.next = self.head.next
        node.prev = self.head
        self.head.next.prev = node
        self.head.next = node

    def get(self, key):
        if key not in self.cache: return -1
        node = self.cache[key]
        self._remove(node)
        self._insert_front(node)

```

```

return node.val

def put(self, key, val):
    if key in self.cache:
        self._remove(self.cache[key])
    node = DLLNode(key, val)
    self.cache[key] = node
    self._insert_front(node)
    if len(self.cache) > self.cap:
        lru = self.tail.prev
        self._remove(lru)
        del self.cache[lru.key]

```

Interview takeaway: LRU = HashMap + DLL with dummy head and tail. The dummy nodes eliminate all edge-case null checks. LFU is significantly harder — use HashMap(key→freq) + HashMap(freq→OrderedDict) + min_freq variable.

Real interview use-case: LRU cache (literal LeetCode 146), LFU cache (LeetCode 460), design Twitter, design browser history.

3.12 Segment Tree & Fenwick Tree (BIT)


What it is:

- **Segment Tree:** Tree where each node stores aggregate (sum/min/max) over a range. Supports range queries and point/range updates in $O(\log n)$.
- **Fenwick Tree (Binary Indexed Tree):** Array-based structure using binary representation for prefix sums. Simpler to implement; less general.

Internal mechanics:

Array: [1, 3, 5, 7, 9, 11] (0-indexed)

Segment Tree (sum):



Query sum(1,4): decompose into [1-2] and [3-4] → 8 + 16 = 24

Update arr[2] = 10: update leaf, propagate to root → $O(\log n)$

Space: $O(4n)$ for array representation

Fenwick Tree (BIT) for prefix sums:

Index:	1	2	3	4	5	6
arr:	[1,	3,	5,	7,	9,11]	
BIT:	[1,	4,	5,16,	9,20]	← each stores sum of specific range	

BIT[i] = sum of arr[i - lowbit(i) + 1 ... i]

lowbit(i) = i & (-i) ← isolates least significant bit

Prefix sum query(5):

i=5 (101): add BIT[5]=9, i=5-(5&5)=4
i=4 (100): add BIT[4]=16, i=4-(4&4)=0 → stop
Total: 9+16=25 ✓ (but wait: 1+3+5+7+9=25 ✓)

Update(3, delta=2): propagate up

i=3: BIT[3]+=2, i=3+(3&3)=4
i=4: BIT[4]+=2, i=4+(4&4)=8 > n → stop

Complexity:

Op	Naive Array	Segment Tree	Fenwick Tree
Build	$O(n)$	$O(n)$	$O(n \log n)$
Point update	$O(1)$	$O(\log n)$	$O(\log n)$
Range update	$O(n)$	$O(\log n)$	$O(\log n)$
Range query	$O(n)$	$O(\log n)$	$O(\log n)$
Space	$O(n)$	$O(4n)$	$O(n)$

When to use:

- › Segment tree: range min/max queries, lazy propagation for range updates
- › Fenwick tree: only prefix sums needed; cleaner code
- › Both: dynamic range queries where array is mutated

Key Python patterns:

```

# Fenwick Tree (BIT) - prefix sum
class BIT:
    def __init__(self, n):
        self.n = n
        self.tree = [0] * (n + 1) # 1-indexed

    def update(self, i, delta): # add delta to position i
        while i <= self.n:
            self.tree[i] += delta
            i += i & (-i) # move to parent

    def query(self, i): # prefix sum [1..i]
        s = 0
        while i > 0:
            s += self.tree[i]
            i -= i & (-i) # move to ancestor
        return s

    def range_query(self, l, r): # sum [l..r]
        return self.query(r) - self.query(l - 1)

# Segment Tree (iterative, cleaner)
class SegmentTree:
    def __init__(self, nums):
        self.n = len(nums)
        self.tree = [0] * (2 * self.n)
        # Build
        for i, v in enumerate(nums):
            self.tree[self.n + i] = v
        for i in range(self.n - 1, 0, -1):
            self.tree[i] = self.tree[2*i] + self.tree[2*i+1]

    def update(self, i, val):
        i += self.n
        self.tree[i] = val
        while i > 1:
            i //= 2
            self.tree[i] = self.tree[2*i] + self.tree[2*i+1]

    def query(self, l, r): # sum [l, r) half-open
        res = 0
        l += self.n; r += self.n
        while l < r:

```

```

if l & 1: res += self.tree[l]; l += 1
if r & 1: r -= 1; res += self.tree[r]
l //= 2; r //= 2
return res

```

Interview takeaway: Fenwick tree is simpler to code and sufficient for prefix sum + point update problems. Use segment tree when you need range updates or range min/max. The "count of smaller numbers after self" problem → BIT is the elegant solution.

Real interview use-case: Range sum query (mutable), count of smaller numbers after self, reverse pairs, falling squares.

3.13 Bloom Filter

What it is: A probabilistic, space-efficient data structure that answers "is X possibly in the set?" or "is X definitely NOT in the set?" — with false positives but no false negatives.

Internal mechanics:

Bloom Filter (m=20 bits, k=3 hash functions):

Initial state:

[0]

Insert "cat" (hash1=3, hash2=7, hash3=15):

[0][0][0][1][0][0][0][1][0][0][0][0][0][0][0][1][0][0][0][0]

↑

↑

↑

Insert "dog" (hash1=3, hash2=10, hash3=17):

[0][0][0][1][0][0][0][1][0][0][1][0][0][0][0][1][0][1][0][0]

Note: bit 3 now set by both "cat" and "dog"!

Query "cat": check bits 3,7,15 → all 1 → POSSIBLY in set ✓

Query "hat": hash1=3, hash2=10, hash3=17 → all 1 → POSSIBLY in set x FALSE POSITIVE!

Query "xyz": hash1=1 → 0 → DEFINITELY NOT in set ✓

False positive probability $\approx (1 - e^{(-kn/m)})^k$

where n=items inserted, m=bit array size, k=hash functions

Key: no false negatives (if bloom says NO, it's definitely NO)

cannot delete items (counter-based bloom filters can)

cannot enumerate stored elements

Complexity:

Op	Time	Space
Insert	O(k)	O(m)
Query	O(k)	O(1)
Delete	N/A	-

When to use:

- › Cache pre-filtering: "has this URL been seen before?" before hitting DB
- › Distributed systems: "does key exist on this node?" before network call
- › Spell checkers: fast "not a word" filter
- › HBase/Cassandra: skip SSTable reads for non-existent keys

Ties to systems:

- › Cassandra uses bloom filters to avoid unnecessary disk I/O
- › Google Chrome used bloom filter for Safe Browsing URL blacklist

- › Bitcoin uses bloom filters in SPV (Simplified Payment Verification)

Interview takeaway: Bloom filters trade accuracy for space and speed. False positives OK; false negatives are impossible. When an interviewer asks "how would you check if a URL was seen before using $O(1)$ space?" — bloom filter is the answer. Note: you can NOT delete from a standard bloom filter.

4 Choosing the Right Data Structure

Decision table: "I need X → use Y"

Need	Use	Why
$O(1)$ access by index	Array / List	Contiguous memory
$O(1)$ key-value lookup	Hash Map	Hash function
Ordered key-value, range queries	BST / Sorted Map	In-order traversal
Get min/max repeatedly	Heap	Root always min/max
Get kth smallest/largest	Heap of size k	$O(n \log k)$
Prefix search / autocomplete	Trie	Shared prefix compression
LIFO (undo, recursion)	Stack	Push/pop from one end
FIFO (BFS, scheduling)	Queue / Deque	Enqueue/dequeue $O(1)$
Sliding window max/min	Monotonic Deque	Maintain monotone order
Next greater/smaller element	Monotonic Stack	Classic pattern
Connected components	Union-Find	Near- $O(1)$ union/find
Shortest path (unweighted)	BFS (Queue)	Level-by-level expansion
Shortest path (weighted, non-neg)	Dijkstra (Heap)	Greedy + priority queue
Shortest path (negative weights)	Bellman-Ford	DP approach
All pairs shortest path	Floyd-Warshall	$O(V^3)$ DP
Range sum query + updates	Fenwick Tree / Segment Tree	$O(\log n)$ both
Range min/max query + updates	Segment Tree	BIT can't do min/max
$O(1)$ LRU eviction	HashMap + DLL	$O(1)$ access + $O(1)$ order
Membership test, space-constrained	Bloom Filter	Probabilistic, tiny space
Count frequencies	HashMap / Counter	Hash map with integer values
Find duplicates	HashSet or sorting	$O(n)$ with hash, $O(n \log n)$ sort
Merge k sorted sequences	Min-Heap of k elements	$O(n \log k)$
Detect cycle (undirected graph)	Union-Find	Simpler than DFS coloring
Detect cycle (directed graph)	DFS with in-stack tracking	Union-Find doesn't work here
Topological ordering	Kahn's (BFS) or DFS	Works on DAGs only

5 Language Notes

5.1 Python

Need	Python	Notes
Dynamic array	<code>list</code>	<code>append</code> $O(1)$ amortized
Hash map	<code>dict</code>	<code>defaultdict(list/int/set)</code>

Need	Python	Notes
Hash set	set	add, in O(1)
Ordered map (by key)	sortedcontainers.SortedList	Not stdlib; use bisect + list
Min heap	heapq (module)	Negate for max heap
Queue / Deque	collections.deque	appendleft, popleft O(1)
Counter	collections.Counter	most_common(k)
LRU built-in	collections.OrderedDict	move_to_end, popitem
Bisect (sorted insert)	bisect.bisect_left/right	Binary search on sorted list
Infinity	float('inf')	
Char to int	ord('a') → 97	chr(97) → 'a'

```

# Common Python idioms
from collections import defaultdict, Counter, deque, OrderedDict
import heapq
from bisect import bisect_left, bisect_right, insert

# Heap with custom key (push tuples)
heapq.heappush(heap, (priority, item))

# Sorted list simulation with bisect
import bisect
sorted_list = []
bisect.insort(sorted_list, 5)      # insert and maintain order
idx = bisect.bisect_left(sorted_list, 5) # binary search

# defaultdict prevents KeyError
graph = defaultdict(list)
count = defaultdict(int)

```

5.2 Java

Need	Java Class	Notes
Dynamic array	ArrayList<T>	add() O(1) amortized
Hash map	HashMap<K,V>	Not ordered
Ordered map (by key)	TreeMap<K,V>	Red-Black tree, O(log n) ops
Ordered set	TreeSet<T>	Red-Black tree
Hash set	HashSet<T>	
Linked hash map	LinkedHashMap<K,V>	Insertion-order iteration (LRU)
Min/Max heap	PriorityQueue<T>	Min by default; pass Comparator.reverseOrder() for max
Stack	Deque<T> as stack	Prefer ArrayDeque over Stack
Queue	ArrayDeque<T>	Or LinkedList<T>
Deque	ArrayDeque<T>	

```

// Java idioms
PriorityQueue<Integer> minHeap = new PriorityQueue<>();
PriorityQueue<Integer> maxHeap = new PriorityQueue<>(Collections.reverseOrder());

Map<String, Integer> map = new HashMap<>();

```

```

map.getOrDefault(key, 0);
map.computeIfAbsent(key, k -> new ArrayList<>()).add(val);

Deque<Integer> stack = new ArrayDeque<>(); // use as stack
stack.push(1); stack.pop(); stack.peek();

Queue<Integer> queue = new ArrayDeque<>(); // use as queue
queue.offer(1); queue.poll(); queue.peek();

TreeMap<Integer, Integer> ordered = new TreeMap<>();
ordered.floorKey(x); // largest key ≤ x
ordered.ceilingKey(x); // smallest key ≥ x

```

5.3 C++

Need	C++ Container	Notes
Dynamic array	<code>vector<T></code>	<code>push_back</code> $O(1)$ amortized
Hash map	<code>unordered_map<K, V></code>	$O(1)$ avg
Ordered map (by key)	<code>map<K, V></code>	Red-Black tree, $O(\log n)$
Hash set	<code>unordered_set<T></code>	
Ordered set	<code>set<T></code>	+ <code>lower_bound</code> , <code>upper_bound</code>
Priority queue (max)	<code>priority_queue<T></code>	Max by default
Priority queue (min)	<code>priority_queue<T, vector<T>, greater<T>></code>	
Stack	<code>stack<T></code>	
Queue	<code>queue<T></code>	
Deque	<code>deque<T></code>	

```

// C++ idioms
#include <bits/stdc++.h>
using namespace std;

unordered_map<int, int> freq;
freq[x]++; // default-initializes to 0

// Min heap
priority_queue<int, vector<int>, greater<int>> minHeap;

// Ordered set with lower_bound
set<int> s;
auto it = s.lower_bound(x); //  $O(\log n)$ 

// Sort with custom comparator
sort(v.begin(), v.end(), [](auto& a, auto& b){ return a[0] < b[0]; });

```

6 Real-Life Analogies

Data Structure	Analogy
Array	Numbered mailboxes in a row — access box 42 instantly
Hash Map	Library card catalog — look up by title, get shelf location
Linked List	Train cars — each car knows only the next car
Stack	Stack of plates — only access the top
Queue	Queue at a bank — first in, first served
Deque	Train that can board/exit from both ends
BST	Phone book — organized so you can binary search
Heap	Priority ER room — most critical patient always treated first
Trie	Autocomplete on your phone — shared prefix saves space
Graph	Road map — cities are nodes, roads are edges
Union-Find	Social groups — merging friend circles
Segment Tree	Population pyramid — sum over any age range quickly
Fenwick Tree	Odometer with clever bit tricks for partial sums
LRU Cache	Small whiteboard — erase what was written longest ago for space
Bloom Filter	Bouncer list — might let someone in wrongly, never turns away regulars

7 Common Interview Questions

7.1 Arrays

- › Two Sum (LeetCode 1)
- › Best Time to Buy and Sell Stock (LC 121)
- › Maximum Subarray / Kadane's (LC 53)
- › Container With Most Water (LC 11)
- › Trapping Rain Water (LC 42)
- › Rotate Array (LC 189)
- › Product of Array Except Self (LC 238)
- › Merge Intervals (LC 56)

7.2 Strings

- › Longest Substring Without Repeating Characters (LC 3)
- › Minimum Window Substring (LC 76)
- › Group Anagrams (LC 49)
- › Valid Palindrome (LC 125)
- › Encode and Decode Strings (LC 271)
- › Longest Palindromic Substring (LC 5)
- › String to Integer (atoi) (LC 8)

7.3 Hash Maps/Sets

- › Two Sum (LC 1)
- › Top K Frequent Elements (LC 347)
- › Longest Consecutive Sequence (LC 128)
- › Subarray Sum Equals K (LC 560)
- › Valid Sudoku (LC 36)
- › Find All Anagrams in a String (LC 438)

7.4 Linked Lists

- › Reverse Linked List (LC 206)
- › Merge Two Sorted Lists (LC 21)
- › Linked List Cycle (LC 141)
- › Reorder List (LC 143)
- › Remove Nth Node From End (LC 19)
- › Merge K Sorted Lists (LC 23)

7.5 Stacks/Queues

- › Valid Parentheses (LC 20)
- › Min Stack (LC 155)
- › Daily Temperatures (LC 739)
- › Largest Rectangle in Histogram (LC 84)
- › Sliding Window Maximum (LC 239)
- › Evaluate Reverse Polish Notation (LC 150)

7.6 Trees

- › Maximum Depth of Binary Tree (LC 104)
- › Invert Binary Tree (LC 226)
- › Same Tree (LC 100)
- › Binary Tree Level Order Traversal (LC 102)
- › Validate Binary Search Tree (LC 98)
- › Lowest Common Ancestor of BST (LC 235)
- › Kth Smallest Element in BST (LC 230)
- › Serialize and Deserialize Binary Tree (LC 297)
- › Diameter of Binary Tree (LC 543)

7.7 Heaps

- › Kth Largest Element in an Array (LC 215)
- › K Closest Points to Origin (LC 973)
- › Find Median from Data Stream (LC 295)
- › Task Scheduler (LC 621)
- › Merge K Sorted Lists (LC 23)

7.8 Tries

- › Implement Trie (LC 208)
- › Word Search II (LC 212)
- › Replace Words (LC 648)
- › Design Search Autocomplete System (LC 642)

7.9 Graphs

- › Number of Islands (LC 200)
- › Clone Graph (LC 133)
- › Pacific Atlantic Water Flow (LC 417)
- › Course Schedule (LC 207) — cycle detection
- › Number of Connected Components (LC 323)
- › Word Ladder (LC 127) — BFS shortest path
- › Network Delay Time (LC 743) — Dijkstra

- › Cheapest Flights Within K Stops (LC 787) — modified Dijkstra/Bellman-Ford

7.10 Union-Find

- › Number of Provinces (LC 547)
- › Redundant Connection (LC 684)
- › Accounts Merge (LC 721)
- › Satisfiability of Equality Equations (LC 990)

7.11 LRU/LFU

- › LRU Cache (LC 146)
- › LFU Cache (LC 460)

7.12 Segment Tree / BIT

- › Range Sum Query — Mutable (LC 307)
- › Count of Smaller Numbers After Self (LC 315)
- › Reverse Pairs (LC 493)

8 Typical Mistakes Candidates Make

1. **Using `list.pop(0)` for queues** — It's $O(n)$. Always use `collections.deque` and `popleft()`.
2. **String concatenation in loops** — `result += s` inside a loop is $O(n^2)$. Use `"".join(parts)`.
3. **Not handling the empty/single-element edge case** — For linked lists, trees, arrays. Check `if not root`, `if not nums`, `if len(nums) < 2` early.
4. **BST validation using only local `left < root < right`** — This misses nodes in subtrees. Pass bounds `(lo, hi)` down the recursion.
5. **Modifying a collection while iterating it** — Causes `RuntimeError`. Copy or collect indices first.
6. **Forgetting to mark nodes visited before enqueueing (BFS)** — Leads to exponential revisits. Mark visited when you enqueue, not when you dequeue.
7. **Using a set as a hash map key** — Sets are unhashable. Use `frozenset` or `tuple(sorted(...))`.
8. **Forgetting Python's `heapq` is min-heap only** — For max-heap, push `-val` and negate on pop.
9. **$O(n^2)$ nested loop when $O(n)$ hash map exists** — When you write two nested loops looking for a pair, always ask: "can I precompute what I'm looking for?"
10. **Not using dummy head node for linked lists** — Head deletion becomes a special case. Dummy head eliminates it.
11. **Recursion without memoization** — Fibonacci, path counting, etc. naturally revisit subproblems. Add `@lru_cache` or a memo dict.
12. **Union-Find without path compression** — Without it, find is $O(n)$. Always compress.
13. **Assuming dict maintains insertion order in Python 2** — It doesn't. Python 3.7+ guarantees it; don't rely on this in cross-version code.
14. **Off-by-one in sliding window** — When window is `[l, r]`, size is `r - l + 1`. Sketch a small example before coding.
15. **Passing arrays by reference and mutating them** — In Python lists are passed by reference. If you need to preserve the original, copy with `arr[:]`.

9 Revision Cheat Sheet

9.1 Complexities to Memorize (Absolute Must)

Array access:	$O(1)$	Hash map get/put:	$O(1)$ avg
Array search:	$O(n)$	BST search:	$O(\log n)$ balanced
Append to list:	$O(1)$ A	Heap push/pop:	$O(\log n)$
List insert/del:	$O(n)$	Build heap:	$O(n)$ ← non-obvious!
Trie ops:	$O(m)$	Segment tree:	$O(\log n)$ query/update
BFS/DFS:	$O(V+E)$	Union-Find:	$O(\alpha(n)) \approx O(1)$
Dijkstra:	$O((V+E) \log V)$		
Heapsort:	$O(n \log n)$		
Sort:	$O(n \log n)$		
Binary search:	$O(\log n)$		
String join:	$O(n)$	String concat loop:	$O(n^2)$ ← DANGER

A = amortized

9.2 Decision Hooks (6 questions to ask every interview)

1. **"Do I need $O(1)$ lookup?"** → Hash map or hash set.
2. **"Do I need sorted order or range queries?"** → BST, sorted array + binary search, or segment tree.
3. **"Do I need repeated min/max?"** → Heap.
4. **"Is this a traversal/reachability problem?"** → BFS (shortest path) or DFS (connectivity, cycles).
5. **"Am I connecting/grouping things?"** → Union-Find.
6. **"Do I need prefix operations?"** → Trie (strings) or Fenwick tree (numbers).

9.3 Quick Pattern Mapping

Pattern	Structure
Sliding window	Array + two pointers
Sliding window max/min	Monotonic deque
Next greater element	Monotonic stack
Level-order traversal	Queue (BFS)
Path/exhaustive search	Stack (DFS) / recursion
Shortest path (unweighted)	BFS
Shortest path (weighted)	Dijkstra (min-heap)
Connected components	BFS/DFS or Union-Find
Top K elements	Heap of size K
Frequency counting	Hash map / Counter
Deduplication	Hash set
Prefix sum	Fenwick tree / prefix array
Range sum (mutable)	Fenwick tree
Range min/max (mutable)	Segment tree
Word prefix search	Trie
Order + $O(1)$ eviction	LRU (HashMap + DLL)
Large-scale membership	Bloom filter

9.4 The 5 Structures You MUST Be Able to Code From Scratch

1. **Hash Map** (open addressing or chaining — at least explain internals)
2. **Linked List** with insert, delete, reverse
3. **Binary Heap** with push, pop, heapify
4. **Trie** with insert, search, startsWith

5. LRU Cache (HashMap + doubly linked list)

9.5 Memory Aids

- › **Heap is NOT a BST** — heap has parent $\leq$ children (only), BST has left $<$ parent $<$ right.
- › **Trie depth = key length, not n** — n = number of keys.
- › **Union-Find has no delete** — it's append-only by design.
- › **BFS finds shortest path; DFS does not** — in unweighted graphs.
- › **heapify is $O(n)$** — this surprises everyone. Proof: most nodes are at the bottom and need to sift down only a little.
- › **Hash map worst case $O(n)$** — when all keys hash to same bucket (adversarial input or bad hash). In practice $O(1)$.
- › **AVL is stricter than Red-Black** — AVL: max height diff 1; RB: height $\leq 2 \cdot \log(n)$. RB faster for insert/delete, AVL faster for lookup.

10 Visual Reference

A set of figures distilling the key quantitative intuitions of this topic.

Data structure operation complexity (average case)				
	Access	Search	Insert	Delete
Array	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Dyn. Array	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Linked List	$O(n)$	$O(n)$	$O(1)$	$O(1)$
Hash Table	$O(1)$	$O(1)$	$O(1)$	$O(1)$
BST (bal.)	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$
Heap	$O(n)$	$O(n)$	$O(\log n)$	$O(\log n)$
Stack/Queue	$O(1)$	$O(n)$	$O(1)$	$O(1)$

Figure 1. Pick the structure that makes your hot operation $O(1)/O(\log n)$. Hash tables win for lookups; balanced trees keep order; heaps give cheap min/max.