Design Track

Interview Preparation Series • Professional Edition

System Design Problems — Foundational

Compiled with LaTeX • June 15, 2026

harsha.somisetty01@gmail.com

In This Edition

1 System Design --- Part 1: Foundational Problems

> Complete end-to-end worked designs for a 45–60 min FAANG system-design round. Each section mirrors the structure you should drive in the actual interview.

1.1 1. Design a URL Shortener (TinyURL / bit.ly)

1. Requirements

Clarifying questions to ask the interviewer:

- What is the expected scale? (DAU, writes per second, reads per second?)
- Should we support custom aliases (user-defined short codes)?
- What is the required URL lifetime? Do links expire?
- Do we need analytics (click counts, geo, referrer)?
- Should we support link editing or deletion?
- Is high availability or strong consistency more important?

Functional Requirements:

1. Given a long URL, return a unique short URL (e.g., `https://tny.io/aB3kX9`).
2. Redirect a short URL to the original long URL.
3. (Optional) Custom aliases: user can specify the short code.
4. (Optional) URL expiry: links expire after N days.
5. (Optional) Analytics: track click count, geography, referrer.

Non-Functional Requirements:

- 100 ms P99 redirect latency.
- 99.99% availability (URL shortener is a critical dependency for customers).
- Short codes must be unique and not predictable/enumerable.
- The system is **read-heavy** (100:1 read-to-write ratio is typical).

2. Capacity Estimation

Assumptions:

- 500 million URLs shortened per month (writes).
- 100:1 read-to-write ratio → 50 billion redirects per month.

Write QPS:

```
500,000,000 / (30 × 86,400) ≈ 193 writes/sec  ≈ ~200 writes/sec
Peak (3×):                                      ≈ ~600 writes/sec
```

Read QPS:

```
50,000,000,000 / (30 × 86,400) ≈ 19,290 reads/sec  ≈ ~20,000 reads/sec
Peak (3×):                                           ≈ ~60,000 reads/sec
```

Storage (URLs, 5 years):

```
500M writes/month × 12 months × 5 years = 30 billion URLs
Each record ≈ 500 bytes (long URL ~400 B + metadata ~100 B)
Total: 30B × 500 B = 15 TB
```

Cache (top 20% URLs serve 80% traffic):

```
20,000 reads/sec × 86,400 sec/day = 1.728 billion reads/day
20% of 30B URLs = 6 billion "hot" URLs → not all fit in RAM
Pragmatic: cache last 24 hrs of writes = 500M/30 × 500B ≈ 8.3 GB/day  →  ~10 GB RAM
```

Bandwidth:

```
Inbound (writes):  200/sec × 500 B  ≈ 100 KB/s
Outbound (reads):  20,000/sec × 500 B ≈ 10 MB/s
```

3. API Design

```
# Shorten a URL
POST /api/v1/shorten
Request:
  {
    "long_url":   "https://www.example.com/very/long/path?q=something",
    "custom_code": "my-alias",          // optional
    "expires_at":  "2027-01-01T00:00Z" // optional
  }
Response 201:
  {
    "short_url": "https://tny.io/aB3kX9",
    "short_code": "aB3kX9",
    "expires_at": "2027-01-01T00:00Z"
  }

# Redirect (the critical read path)
GET /{short_code}
Response 302:
  Location: https://www.example.com/very/long/path?q=something

# Analytics (optional)
GET /api/v1/stats/{short_code}
Response 200:
  {
    "clicks": 14982,
    "top_countries": [...],
    "referrers": [...]
  }

# Delete
DELETE /api/v1/shorten/{short_code}
Response 204
```

4. Data Model

SQL vs NoSQL decision:

- The primary access pattern is a simple key-value lookup: `short_code → long_url`.
- No complex joins needed; writes are straightforward inserts.
- **Choice: NoSQL key-value store (Cassandra or DynamoDB) for the URL table** — horizontal scale, high availability, and the access pattern is a perfect fit.
- A relational DB (PostgreSQL) is used for user/account data where consistency matters.

URL Table (Cassandra / DynamoDB):

```
urls

short_code    TEXT       PRIMARY KEY        -- "aB3kX9"
long_url      TEXT       NOT NULL           -- original URL
user_id       UUID                          -- creator (nullable for anonymous)
created_at    TIMESTAMP
expires_at    TIMESTAMP                     -- NULL = never expires
is_deleted    BOOLEAN    DEFAULT false
```

Users Table (PostgreSQL):

```
users

user_id       UUID       PRIMARY KEY
email         VARCHAR(255) UNIQUE NOT NULL
api_key       VARCHAR(64)  UNIQUE            -- for rate limiting
created_at    TIMESTAMP
```

Analytics Table (write-optimized, e.g., Cassandra or ClickHouse):

```
click_events

short_code    TEXT
clicked_at    TIMESTAMP
country       TEXT
referrer      TEXT
user_agent    TEXT
-- Partition key: short_code; Clustering key: clicked_at DESC
```

5. High-Level Architecture

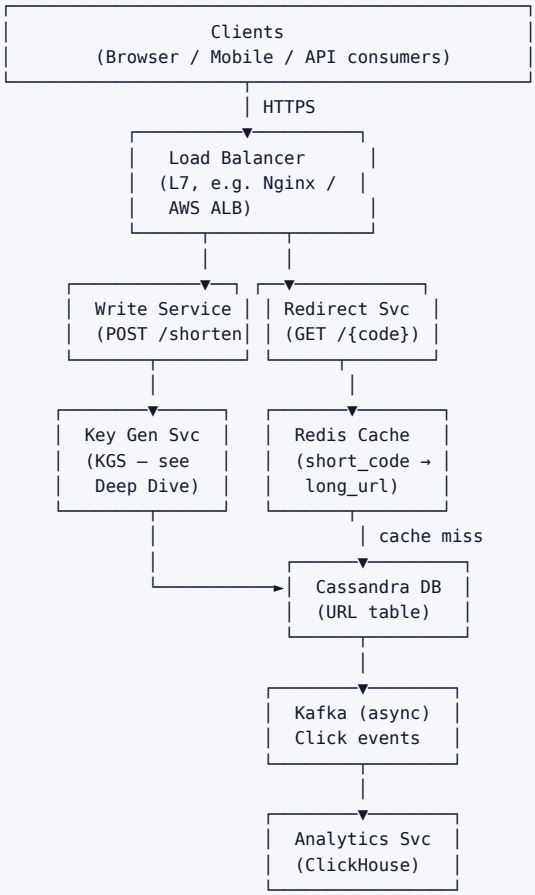

Component walkthrough:

1. **Client** sends `POST /shorten` or `GET /{code}` over HTTPS.
2. **Load Balancer** routes to stateless Write Service or Redirect Service pods.
3. **Redirect Service** checks **Redis** first (cache hit → 302 immediately). On cache miss, reads Cassandra, populates cache, then redirects.
4. **Write Service** calls **KGS** (Key Generation Service) for a pre-generated unique code, stores `(code,`

`long_url)` in Cassandra, and warms the cache.
5. **Kafka** asynchronously receives click events; **Analytics Service** batch-writes to ClickHouse for reporting.

6. Deep Dives

6.1 Key Generation: How to create short codes?

Approach	How	Pros	Cons
MD5 / SHA-256 hash	Hash long URL, take first 6–8 chars	Deterministic (same URL → same code)	Collisions possible; must check DB on every write
Counter + Base62	Global atomic counter → encode in Base62	No collisions; simple	Single counter = SPOF; reveals ordering/volume
Pre-generated KGS	Dedicated service pre-generates & stores unused keys in a ”keys” table	No collision checking at write time; fast	Additional service to maintain; key theft if service crashes
UUID	Generate random UUID, take first 8 chars	Simple	High collision probability at scale

Key Decision: KGS (Key Generation Service) is the recommended approach for large scale.
KGS design:

- KGS pre-generates 6-character Base62 codes (62^6 ≈ 56 billion unique codes).
- Stores them in two tables: `unused_keys` and `used_keys`.
- Each Write Service instance requests a batch (e.g., 1,000 keys) from KGS on startup, holds them in-memory.
- KGS marks those keys as ”in use” atomically. If a Write Service pod crashes, those keys are wasted but never reused — acceptable loss.
- KGS itself can be replicated with a leader for key assignment.

6.2 Redirect: 301 vs 302?

	301 Permanent	302 Temporary
Browser caches?	Yes — browser won't call TinyURL again	No — browser calls TinyURL on every click
Analytics accuracy	**Broken** — browser bypasses your servers	**Accurate** — every click hits your servers
Load on TinyURL	Lower after first visit	Higher (every click)
Recommendation	Use only if analytics not needed	**Use 302 for full analytics capability**

Interview takeaway: Always use 302 unless you explicitly trade analytics for reduced server load.

6.3 Collision Handling If using hash-based approach:

1. Compute MD5(long_url) → take first 7 chars as short_code.
2. Check if `short_code` exists in DB.
3. If collision: append a salt (e.g., increment integer), re-hash, retry.
4. Limit retries to 3; fall back to random generation if exhausted.

With KGS, collisions are impossible — keys are pre-assigned and marked used atomically.

6.4 Custom Aliases

- User provides desired short code (e.g., `my-brand`).
- Check availability: `SELECT 1 FROM urls WHERE short_code = 'my-brand'`.
- If taken → return 409 Conflict.

- Store with `is_custom = true` flag to distinguish from generated codes.
- Rate-limit custom alias creation per user to prevent squatting.

6.5 URL Expiry

- Store `expires_at` timestamp in the URL record.
- On redirect: if `expires_at < NOW()` → return 404 (or 410 Gone).
- Background cleanup job: delete expired records weekly (keep DB lean).
- Cache entries: set Redis TTL = `expires_at - NOW()` so expired URLs auto-evict.

7. Bottlenecks & Scaling

Read hotspot — viral URLs:

- A single URL (e.g., link from a viral tweet) can receive millions of reads/sec.
- Redis handles this well — single key lookup is O(1) and Redis handles 100k+ ops/sec per node.
- For extreme cases: replicate the specific key across multiple Redis nodes or use local in-process cache (LRU with 1000 entries) in each Redirect Service pod.

Write scaling:

- Writes at ~200/sec are trivially handled. Cassandra scales horizontally; add nodes as needed.
- KGS: single KGS with a read replica for HA is sufficient at this scale. Key generation is CPU-light.

DB sharding:

- Cassandra partitions by `short_code` automatically (consistent hashing). No manual sharding needed.
- If using SQL: shard by `short_code` modulo N (hash partitioning). Use a consistent hash ring to avoid resharding pain.

Cache eviction:

- Redis with LRU policy. Set max memory = 10–20 GB.
- TTL on each entry = min(expires_at, 24 hrs) to prevent stale data.

Failure handling:

- Redis down → fall through to Cassandra (latency increases but system stays up).
- KGS down → Write Service uses its in-memory key batch until KGS recovers; alert if batch exhausted.
- Cassandra node down → Cassandra replication factor 3 (RF=3) ensures reads/writes still succeed with 2 nodes up.

Geographic distribution:

- Deploy read replicas of Cassandra in multiple regions.
- Use GeoDNS to route users to nearest Redirect Service cluster.
- URL writes go to a primary region; replicated asynchronously to others.

8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Key generation	**KGS (pre-generated)**	Hash + collision check	No collision risk; predictable latency
DB for URLs	**Cassandra (NoSQL)**	PostgreSQL	Write scalability; simple key-value access pattern
Redirect type	**302 Temporary**	301 Permanent	Enables accurate analytics
Cache	**Redis (LRU)**	Memcached	Rich data types; TTL support; persistence

Decision	Choice Made	Alternative	Reason
Analytics pipeline	**Kafka → ClickHouse (async)**	Sync write to SQL	Decouples write path; handles bursty click events
Custom aliases	**Supported (rate-limited)**	Not supported	Product differentiator; must prevent abuse

1.2 2. Design a Rate Limiter

1. Requirements

Clarifying questions to ask the interviewer:

- Where should the rate limiter live? Client-side, API gateway, or per-service middleware?
- What are we rate-limiting on? Per user? Per IP? Per API key? Per endpoint?
- What happens when a limit is exceeded — hard reject (429) or soft throttle (queue)?
- Do we need distributed rate limiting (multiple servers sharing state)?
- What are the latency requirements? Rate limit check must not add significant overhead.
- Do rules need to be configurable at runtime without deployments?

Functional Requirements:

1. Limit the number of requests a client (user/IP/API key) can make in a time window.
2. Return HTTP 429 Too Many Requests when limit is exceeded, with `Retry-After` header.
3. Support multiple limits (e.g., 100 req/min AND 1,000 req/hr for same client).
4. Rules are configurable without service restarts.

Non-Functional Requirements:

- Rate limit check must add < 5 ms P99 latency overhead.
- Highly available — rate limiter failure should fail open (allow traffic) not fail closed (block everything).
- Works across multiple servers (distributed, consistent limits).
- Accurate enough — exact precision not required; small over-limit acceptable.

2. Capacity Estimation

Assumptions:

- 10 million DAU, each making ~100 requests/day → 1 billion requests/day.
- Rate limit state must be checked for every request.

QPS to rate limiter:

```
1,000,000,000 / 86,400 ≈ 11,574 req/sec  ≈ ~12,000 checks/sec
Peak (3×):                               ≈ ~36,000 checks/sec
```

Memory for rate limit counters:

```
Strategy: sliding window counter per user per endpoint, 1-hr window.
10M users × 10 endpoints × (counter + timestamp) ≈ 10M × 10 × 20 bytes = 2 GB
→ Fits comfortably in a Redis cluster.
```

Bandwidth:

- Rate limit check: small — just a counter lookup and increment.
- Each Redis call ≈ microseconds; no large data transfer.

3. API Design

The rate limiter is typically middleware, not a standalone API. But if exposing as a service:

```
# Check and record a request (called by API gateway)
POST /rate-limit/check
Request:
  {
    "client_id": "user:1234",
    "endpoint":  "/api/v1/posts",
    "timestamp": 1718000000
  }
Response 200 (allowed):
  {
    "allowed": true,
    "remaining": 42,
    "reset_at": 1718003600
  }
Response 429 (throttled):
  {
    "allowed": false,
    "retry_after": 58,       // seconds until window resets
    "limit": 100,
    "window": "1m"
  }

# Update rules (admin)
PUT /rate-limit/rules/{client_id}
Request:
  {
    "rules": [
      { "window": "1m", "limit": 100 },
      { "window": "1h", "limit": 1000 }
    ]
  }
```

Response headers to include on every API response:

```
X-RateLimit-Limit:     100
X-RateLimit-Remaining: 42
X-RateLimit-Reset:     1718003600
Retry-After:           58          // only on 429
```

4. Data Model

Redis data structures (no SQL needed — state is ephemeral):

Token Bucket / Fixed Window counter:

```
Key:   "rl:{client_id}:{endpoint}:{window_start}"
Value: integer (request count)
TTL:   window duration (e.g., 60s for 1-min window)
```

Sliding Window Log:

```
Key:   "rl:{client_id}:{endpoint}"
Type:  Redis Sorted Set
Member: request_id (UUID)
Score:  timestamp (epoch ms)
TTL:   window duration
```

Rules config (read rarely, cache in-process):

```
# Redis Hash or a config store (Consul / etcd)
Key:   "rule:{client_id}"
Value: JSON { "rules": [{"window":"1m","limit":100}, ...] }
```

5. High-Level Architecture

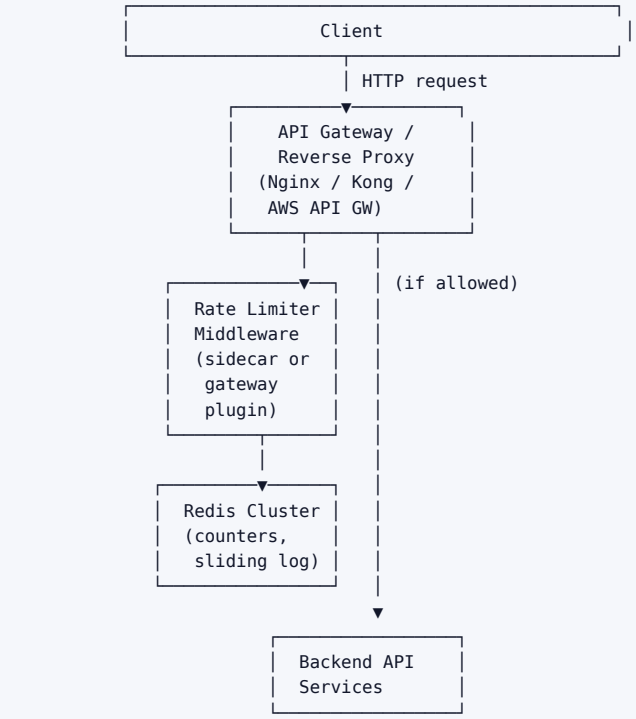

Component walkthrough:

1. Every request hits the **API Gateway**.
2. Gateway invokes **Rate Limiter Middleware** (same process or sidecar).
3. Middleware reads the rule for `(client_id, endpoint)` from its **in-process cache** (refreshed every 30s from etcd/Consul).
4. Middleware atomically increments/checks the counter in **Redis**.
5. If allowed: forward request to **Backend Service**.
6. If denied: return 429 immediately with `Retry-After`.

6. Deep Dives

6.1 Algorithm Comparison

Algorithm	How it works	Pros	Cons	Best for
Fixed Window Counter	Reset counter every N seconds	Simple; O(1) memory	Boundary burst: 2× limit possible at window edge	Simple internal limits
Sliding Window Log	Store timestamp of every request in sorted set; count within window	Accurate	O(n) memory per user (n = requests in window)	Low-volume, high-accuracy needs
Sliding Window Counter	Blend current + previous window counts weighted by overlap	Near-accurate, O(1) memory	Slight inaccuracy (~0.003% over-limit)	**Best general choice**

Algorithm	How it works	Pros	Cons	Best for
Token Bucket	Bucket of tokens refilled at rate R; request consumes one	Allows bursts up to bucket size	Harder to reason about exact limits; burst may confuse users	APIs that allow bursting
Leaky Bucket	Queue requests; process at constant rate	Smooth output rate	Queue introduces latency; bursty requests get delayed	Network traffic shaping

Key Decision: Sliding Window Counter for most API rate limiting. Token Bucket when you want to allow short bursts.

6.2 Sliding Window Counter — Implementation Detail

```
Window = 1 minute.
Current minute bucket: starts at :00
Previous minute bucket: started at :-60

count = prev_bucket_count × (overlap / window) + curr_bucket_count

Example at :45 into the current minute:
overlap = 60 - 45 = 15 seconds (15/60 = 25% of previous window still counts)
count = prev_count × 0.25 + curr_count

If count > limit → reject.
```

Redis implementation:

CODE

```lua
-- Lua script (atomic)
local key_curr = KEYS[1]  -- "rl:user1:1718000060"
local key_prev = KEYS[2]  -- "rl:user1:1718000000"
local limit    = tonumber(ARGV[1])
local now      = tonumber(ARGV[2])
local window   = tonumber(ARGV[3])
local window_start = now - (now % window)
local overlap  = window - (now - window_start)

local curr_count = tonumber(redis.call('GET', key_curr) or 0)
local prev_count = tonumber(redis.call('GET', key_prev) or 0)
local weighted   = prev_count * (overlap / window) + curr_count

if weighted >= limit then
  return 0  -- rejected
end
redis.call('INCR', key_curr)
redis.call('EXPIRE', key_curr, window * 2)
return 1  -- allowed
```

6.3 Where to Place the Rate Limiter?

Placement	Pros	Cons
Client-side	Zero server load	Easily bypassed; untrustworthy
API Gateway	Single enforcement point; language-agnostic	Gateway becomes a bottleneck; less context
Service middleware	Per-service granularity; business context	Duplicated logic across services; harder to change rules centrally
Dedicated Rate Limit Service	Centralized rules; all services call it	Extra network hop; additional failure point

Placement	Pros	Cons
Recommended: Gateway + Redis	Rules centralized in Redis; enforced at gateway	Need to handle Redis failure gracefully

6.4 Distributed Rate Limiting — Race Conditions **Problem:** Two requests arrive simultaneously on different servers. Both read count=99 (limit=100), both increment to 100, both allow. Result: 2 requests slip through over limit.

Solutions:

1. **Redis Lua scripts (recommended):** Atomic read-increment-check in a single Lua script. Redis is single-threaded; no race condition possible.
2. **Redis INCR + EXPIRE pattern:**

```
count = INCR(key)
if count == 1: EXPIRE(key, window)  -- set TTL only on first request
if count > limit: return 429
```

Problem: race between INCR and EXPIRE if server crashes. Mitigate with `SET key 0 EX window NX` on first call.

3. **Sticky sessions / consistent hashing:** Route same client_id to same rate limiter node. Eliminates distributed counting but adds routing complexity and uneven load.

Key Decision: Use Redis Lua scripts for atomicity. No locks, no transactions overhead.

6.5 Fail-Open vs Fail-Closed

- **Fail-open (recommended):** If Redis is down, allow all requests. Availability > perfect rate limiting.
- **Fail-closed:** If Redis is down, reject all requests. Too severe for most use cases — takes down your API.
- Implementation: wrap Redis call in try-catch; on exception, log and allow the request, alert on-call.

7. Bottlenecks & Scaling

Redis as single point of failure:

- Use Redis Cluster (sharded) — each client_id maps to a shard via consistent hashing.
- Redis Sentinel for HA (automatic failover).
- Replication lag between primary and replica: acceptable (small over-limit possible).

Redis latency:

- Redis P99 ≈ 0.5–1 ms in-datacenter. Total rate limit overhead < 2 ms.
- If Redis is co-located on same host (Redis sidecar per gateway node): P99 < 0.2 ms.

Rule changes at scale:

- Store rules in etcd/Consul. Rate limiter middleware subscribes to watch-notifications.
- In-process cache with 30s TTL: small window where old rules apply after update.

Hotspots (one user generating massive traffic):

- Rate limiter itself handles this well — a single Redis key gets hammered, but Redis is designed for this.
- Use local in-process counter as first layer (no Redis call for clearly non-bursting clients), then Redis for enforcement.

8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Algorithm	**Sliding Window Counter**	Token Bucket	Near-exact accuracy; O(1) memory; prevents boundary bursts
State store	**Redis Cluster**	In-process memory	Distributed consistency across servers; fast enough
Placement	**API Gateway middleware**	Dedicated service	Avoids extra network hop; centralizes enforcement
Atomicity	**Redis Lua scripts**	Redis WATCH/MULTI	Simpler; no retry logic needed
Failure mode	**Fail-open**	Fail-closed	Availability > perfect accuracy for most systems
Rules storage	**etcd + in-process cache**	DB query per request	Rules rarely change; in-process cache eliminates lookup overhead

1.3 3. Design a News Feed (Facebook/Twitter Timeline)

1. Requirements

Clarifying questions to ask the interviewer:

- Is this a ”friends” model (Facebook — bidirectional) or ”followers” model (Twitter — unidirectional)?
- What is the feed ordering — chronological or ranked by relevance?
- What's the expected social graph size? Max followers per user?
- Do we support stories (ephemeral) vs. posts (permanent)?
- Is the feed pre-generated (push) or assembled on request (pull)?
- Do we need to support ads injection in the feed?

Functional Requirements:

1. User can create a post (text, image, video link).
2. User sees a feed of posts from people they follow, ordered by recency or ranking.
3. Feed updates in near-real-time (within seconds of post creation).
4. Pagination: load N posts, scroll for more.
5. Users can like, comment, share (out of scope for this design; focus on feed generation).

Non-Functional Requirements:

- 500 million DAU (Facebook scale).
- Feed load latency < 200 ms.
- Post creation propagates to followers within 5 seconds (eventual consistency is OK).
- Users can have up to 5,000 friends (Facebook) or millions of followers (Twitter celebrities).

2. Capacity Estimation

Assumptions:

- 500 million DAU.
- Each user views feed 5 times/day → 2.5 billion feed reads/day.
- Each user creates 1 post every 5 days → 100 million new posts/day.
- Average friends/followers: 300.

Post write QPS:

```
100,000,000 / 86,400 ≈ 1,157 writes/sec  ≈ ~1,200 posts/sec
Peak (5×):                               ≈ ~6,000 posts/sec
```

Feed read QPS:

```
2,500,000,000 / 86,400 ≈ 28,935 reads/sec  ≈ ~29,000 reads/sec
Peak (3×):                                  ≈ ~87,000 reads/sec
```

Fanout writes (push model — writing to all follower feeds on post):

```
1,200 posts/sec × 300 avg followers = 360,000 feed writes/sec
Celebrity (10M followers): 1 post → 10M writes (burst)
```

Storage:

```
Posts: 100M/day × 1 KB (text) = 100 GB/day text
       Images/videos: separate blob store (S3)
Feed cache: 500M users × 20 cached posts × 1 KB = 10 TB → too big for RAM
  → Cache only active users (10% daily active at peak):
    50M users × 20 posts × 1 KB = 1 TB → feasible with Redis cluster
```

3. API Design

```
# Create a post
POST /api/v1/posts
Auth: Bearer token
Request:
  {
    "content":    "Hello world!",
    "media_urls": ["https://cdn.example.com/img/abc.jpg"],
    "visibility": "friends"   // or "public", "private"
  }
Response 201:
  {
    "post_id":    "p_9x3kA2",
    "created_at": "2026-06-14T10:00:00Z"
  }

# Get news feed
GET /api/v1/feed?limit=20&cursor=<opaque_cursor>
Auth: Bearer token
Response 200:
  {
    "posts": [
      {
        "post_id":    "p_9x3kA2",
        "author_id":  "u_4f2k",
        "author_name":"Alice",
        "content":    "Hello world!",
        "media_urls": [...],
        "like_count": 142,
        "comment_count": 7,
        "created_at": "2026-06-14T10:00:00Z"
      },
      ...
    ],
    "next_cursor": "<opaque_cursor>"
  }

# Get a user's profile posts
GET /api/v1/users/{user_id}/posts?limit=20&cursor=<cursor>
Response 200: { "posts": [...], "next_cursor": "..." }
```

4. Data Model

SQL vs NoSQL:

- **Posts table:** Write-once, read-many. High read QPS. Shard by `author_id`. → NoSQL (Cassandra) or sharded PostgreSQL.
- **Social graph (follows):** Relationship lookups (”who does user X follow?”). → Graph DB (Neo4j) or sharded SQL; at Facebook scale: custom TAO (graph store).

- **Feed cache:** Redis sorted sets (score = post timestamp).
- **User data:** Relational (PostgreSQL) — low write volume, complex queries.

Posts Table (Cassandra — partition by author, cluster by time):

```
posts

author_id   UUID        PARTITION KEY
post_id     TIMEUUID    CLUSTERING KEY (DESC)   -- newest first
content     TEXT
media_urls  LIST<TEXT>
like_count  COUNTER
created_at  TIMESTAMP
```

Social Graph Table (Cassandra — two tables for O(1) lookups):

```
follows_by_follower          -- "who does user X follow?"

follower_id  UUID  PARTITION KEY
followee_id  UUID  CLUSTERING KEY

follows_by_followee          -- "who follows user X?"

followee_id  UUID  PARTITION KEY
follower_id  UUID  CLUSTERING KEY
```

Feed Cache (Redis Sorted Set per user):

```
Key:     "feed:{user_id}"
Members: post_id
Scores:  timestamp (epoch ms) — higher = newer
TTL:     48 hours (evict inactive users)
Max size: 800 posts per user (trim older entries)
```

5. High-Level Architecture

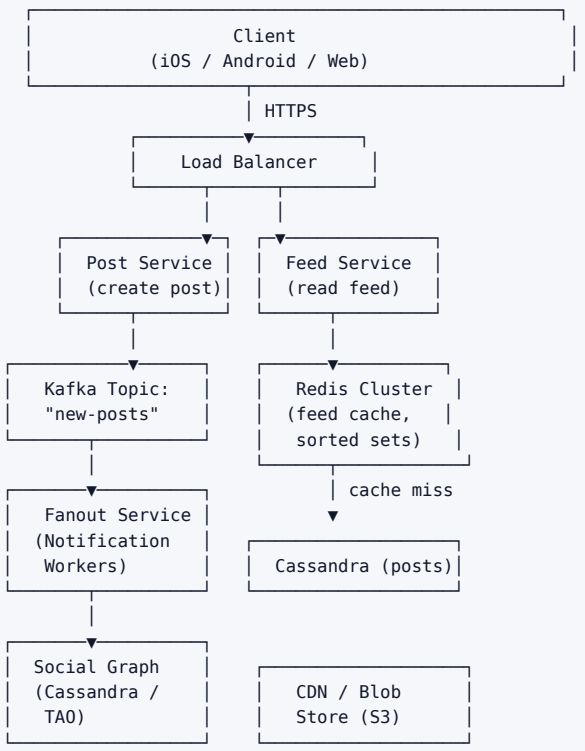

```
Post creation flow:  Post Svc → Cassandra + Kafka
Fanout flow:         Kafka → Fanout Workers → Redis (write to follower feeds)
Feed read flow:      Feed Svc → Redis cache → (miss) Cassandra
```

6. Deep Dives

6.1 Fanout Strategies

Strategy	How	Pros	Cons	Best for
Fanout-on-Write (Push)	On post creation, write post_id to each follower's feed cache	Feed reads O(1) — just read from cache	Write amplification: 1 post → N writes; celebrity (10M followers) = 10M cache writes	Average users
Fanout-on-Read (Pull)	On feed request, fetch posts from each followee, merge, sort	No write amplification	Feed read O(N followees) — slow for users following many people	Celebrity accounts, inactive followers
Hybrid (recommended)	Push to regular followers; pull for celebrities at read time; don't push to inactive users	Balances read and write load	More complex logic	**Large-scale production systems**

Key Decision: Hybrid fanout.
Implementation:

- Define ”celebrity” as users with > 1M followers.
- On post write: fanout-on-write for non-celebrities; for celebrities, mark the post in a ”celebrity posts” index only.
- On feed read: read user's pre-built feed from Redis (push-mode posts) + fetch latest celebrity posts from celebrity index + merge and deduplicate.
- Skip fanout for users who haven't been active in 7 days (check last_active timestamp).

6.2 Celebrity Problem (Hotspot) A celebrity with 10 million followers posts → need 10 million cache writes within seconds.

Mitigation:

1. **Async fanout workers:** Put the fanout job on Kafka with 10,000 partitions. Workers consume in parallel — 10M writes spread across workers.
2. **Pull for celebrity content:** Don't write celebrity posts to follower feeds. At read time, the top-K celebrities a user follows are fetched separately and injected into the feed.
3. **Pre-fetch:** For celebrities, replicate their latest posts to a separate, highly-cached ”celebrity feed” shard. All followers read from the same shard (CDN-like effect).

6.3 Feed Ranking Simple timeline: sort by `score = timestamp`.
ML-ranked feed (Facebook/Instagram model):

```
score = f(recency, engagement_velocity, relationship_strength, content_type)

engagement_velocity = likes_per_hour / avg_likes_per_hour_for_author
relationship_strength = interaction_history(viewer, author)
```

Interview takeaway: In the interview, acknowledge ranking but defer details. Say: ”We'd train an ML model offline, serve predictions via a Feature Store + low-latency inference service. The feed service re-ranks the top 200 candidates before returning the top 20.”

6.4 Pagination Cursor-based pagination (not offset):

- Offset pagination (`LIMIT 20 OFFSET 100`) breaks when new posts are inserted — items shift, causing duplicates/skips.
- Cursor = encoded last-seen post timestamp + post_id.
- Redis: `ZRANGEBYSCORE feed:{user_id} -inf <cursor_score> LIMIT 0 20`.
- Cursor makes the feed stable even as new posts arrive.

7. Bottlenecks & Scaling

Feed cache size:

- 10 TB for active users — run 50+ Redis nodes (200 GB each) in a cluster.
- Use consistent hashing for user → Redis shard mapping.
- Evict feeds of users inactive for 7 days to save memory.

Cassandra post storage:

- Partition by `author_id` → hot partition if one author posts very frequently.
- Mitigate: add `bucket` to partition key: `(author_id, bucket)` where `bucket = created_at / 30 days`.

Social graph at scale:

- Facebook built TAO (The Associations and Objects) — a custom graph store on top of MySQL + Memcache.
- At smaller scale: Cassandra `follows_by_follower` + `follows_by_followee` tables work well.

Read path latency breakdown (target < 200 ms):

```
Redis feed read:          5 ms
Celebrity post fetch:    20 ms (parallel)
User profile hydration:  10 ms (from user cache)
Ranking re-score:        15 ms
Network + serialization: 10 ms
Total:                  ~60 ms  (well within 200 ms budget)
```

8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Fanout strategy	**Hybrid (push+pull)**	Pure push or pure pull	Balances write amplification vs read latency
Feed storage	**Redis sorted sets**	DB query at read time	Sub-10ms reads; DB can't handle 87K reads/sec directly
Feed ordering	**Reverse-chronological (MVP)**	ML-ranked feed	Simpler; ranked feed is a layer on top
Pagination	**Cursor-based**	Offset-based	Stable under inserts; no duplicate/skip on new posts
Post DB	**Cassandra**	PostgreSQL (sharded)	Write throughput; natural time-series clustering
Celebrity handling	**Pull at read time**	Fanout to all followers	Avoids 10M+ write storms per celebrity post

1.4 4. Design a Chat System (WhatsApp / Messenger)

1. Requirements

Clarifying questions to ask the interviewer:

- 1:1 chat only, or also group chat? What is the max group size?

- Do we need real-time delivery (online) and push notifications (offline)?
- Media support (images, video, voice)?
- Do messages need end-to-end encryption?
- Message history: how far back? Searchable?
- Read receipts (sent / delivered / read)?
- Online presence (”last seen”)?

Functional Requirements:

1. 1:1 messaging with real-time delivery.
2. Group messaging (up to 500 members per group).
3. Online/offline status with ”last seen” timestamp.
4. Read receipts: sent → delivered → read.
5. Push notifications for offline users.
6. Message history persisted; searchable.
7. Media: images, files (< 100 MB each).

Non-Functional Requirements:

- 500 million DAU; 100 billion messages/day.
- Message delivery latency < 500 ms (online users).
- High availability: 99.99%.
- Messages delivered exactly once and in order.
- End-to-end encryption support (Signal Protocol) — note architecture impact.

2. Capacity Estimation

Assumptions:

- 500 million DAU.
- Each user sends 40 messages/day → 20 billion messages/day (1:1 + group).
- Average message size: 100 bytes (text).
- Average group has 10 members → group message creates 9 additional deliveries.
- Media: 5% of messages are media → 1 billion media messages/day.

Message write QPS:

```
20,000,000,000 / 86,400 ≈ 231,481 msg/sec  ≈ ~230,000 messages/sec
Peak (3×):                                  ≈ ~700,000 messages/sec
```

Storage:

```
Text messages:  20B/day × 100 B         = 2 TB/day
Media:          1B/day × 1 MB avg       = 1 PB/day  → S3/GCS
Metadata:       20B/day × 50 B          = 1 TB/day
Total text+meta ≈ 3 TB/day → 1.1 PB/year (text only)
```

Concurrent WebSocket connections:

```
500M DAU × 30% online simultaneously = 150M concurrent connections
Each connection: 65KB OS overhead
Total: 150M × 65KB ≈ 9.75 TB RAM just for connections
→ Need ~100,000 chat servers @ 1,500 connections each
(In practice: use event-driven servers; Nginx/Go handles 100K connections per server)
→ 1,500 servers (100K connections each)
```

3. API Design

```
# WebSocket connection (real-time channel)
ws://chat.example.com/ws?token=<auth_token>

# WebSocket message format (client → server)
{
  "type":      "message",
  "msg_id":    "client-generated-uuid",   // for deduplication
  "to":        "user:5678",               // or "group:9999"
  "content":   "Hey!",
  "media_url": null,
  "timestamp": 1718000000000             // client time ms
}

# WebSocket message format (server → client, push delivery)
{
  "type":        "message",
  "msg_id":      "server-assigned-uuid",
  "from":        "user:1234",
  "to":          "user:5678",
  "content":     "Hey!",
  "server_ts":   1718000000050,          // server timestamp (canonical ordering)
  "seq":         100034                  // sequence number in conversation
}

# ACK (client → server, delivery acknowledgement)
{ "type": "ack", "msg_id": "...", "ack_type": "delivered" }
{ "type": "ack", "msg_id": "...", "ack_type": "read" }

# REST: Send message (fallback for HTTP clients)
POST /api/v1/messages
Request:  { "to": "user:5678", "content": "Hey!" }
Response: { "msg_id": "...", "server_ts": 1718000000050 }

# REST: Fetch message history
GET /api/v1/conversations/{conversation_id}/messages?before=<msg_id>&limit=50
Response: { "messages": [...], "has_more": true }

# REST: Upload media
POST /api/v1/media/upload
Response: { "media_url": "https://cdn.example.com/m/abc.jpg", "expires_in": 86400 }

# REST: Get online status
GET /api/v1/users/{user_id}/presence
Response: { "online": false, "last_seen": "2026-06-14T09:55:00Z" }
```

4. Data Model

Storage choices:

- **Messages:** HBase / Cassandra — write-heavy time-series; partition by conversation, cluster by time.
- **Conversations/Groups:** Cassandra or DynamoDB.
- **User data:** PostgreSQL (low write, complex queries).
- **Presence (online status):** Redis (ephemeral, TTL-based).
- **Message queue for offline delivery:** Kafka or internal queues.
- **Media:** S3 / GCS (blob).

Messages Table (HBase / Cassandra):

```
messages

conversation_id   TEXT       PARTITION KEY    -- "conv:{user1}_{user2}" or "group:{id}"
seq_num           BIGINT     CLUSTERING KEY DESC  -- monotonically increasing per conversation
msg_id            UUID       -- globally unique
sender_id         UUID
content           TEXT
media_url         TEXT
```

```
msg_type          TEXT       -- "text", "image", "video", "system"
status            TEXT       -- "sent", "delivered", "read"
created_at        TIMESTAMP
```

Conversations Table:

```
conversations

conversation_id   TEXT    PRIMARY KEY
participants      LIST<UUID>
group_name        TEXT    (null for 1:1)
last_msg_id       UUID
last_msg_at       TIMESTAMP
```

User Conversations Index (reverse lookup):

```
user_conversations   -- "what conversations is user X in?"

user_id              UUID    PARTITION KEY
conversation_id      TEXT    CLUSTERING KEY
last_msg_at          TIMESTAMP  (for sorting)
unread_count         INT
```

Presence Table (Redis):

```
Key:   "presence:{user_id}"
Value: { "online": true, "last_heartbeat": 1718000000000 }
TTL:   30 seconds (refreshed by heartbeat every 10s)
```

5. High-Level Architecture

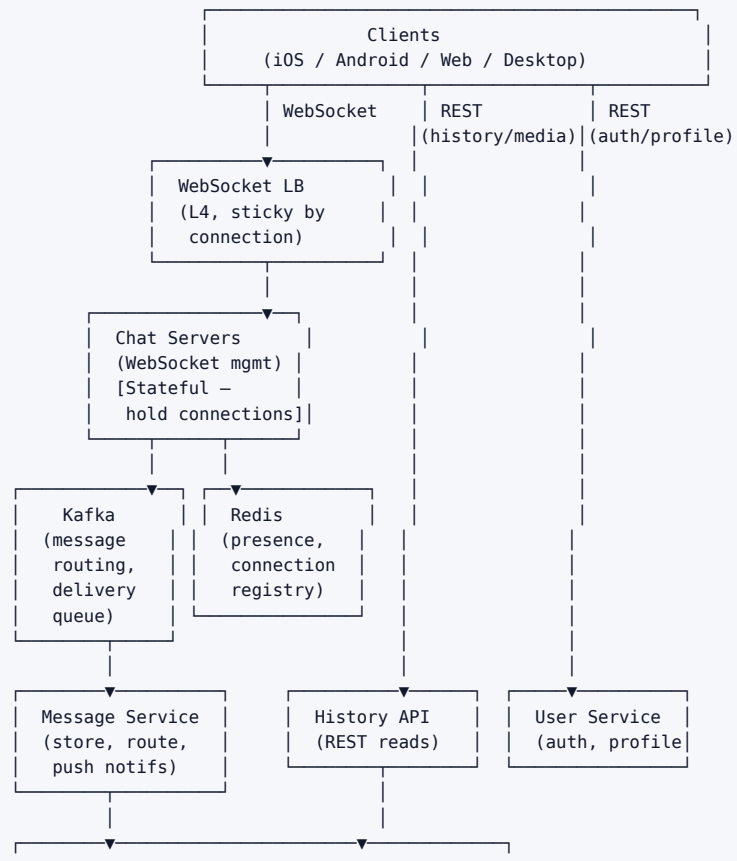

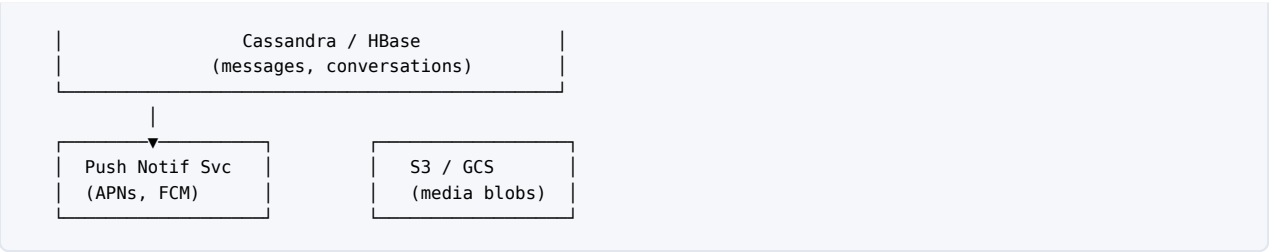

Message delivery flow (1:1, both online):

1. Sender client sends message over WebSocket to **Chat Server A**.
2. Chat Server A writes message to **Kafka** (topic: `messages`).
3. **Message Service** consumes from Kafka, persists to Cassandra.
4. Message Service looks up which **Chat Server** the recipient is connected to (via Redis connection registry: `"ws_server:{user_id}" → "chat-server-B"`).
5. Message Service pushes message to **Chat Server B** (via internal pub/sub).
6. Chat Server B delivers to recipient's WebSocket.
7. Recipient client sends `ack: delivered` → Chat Server B → Message Service → Cassandra (update status) → Kafka → Chat Server A → sender gets double-tick.

6. Deep Dives

6.1 WebSocket Connection Management Why WebSocket instead of HTTP polling?

Method	How	Pros	Cons
Short polling	Client polls every N seconds	Simple	Wasted requests; high latency (up to N sec)
Long polling	Server holds request until message or timeout	Fewer round trips	Still HTTP overhead; server holds threads
WebSocket	Full-duplex TCP connection	True real-time; low overhead	Stateful — servers must track who is connected where
SSE (Server-Sent Events)	Server pushes, client uses XHR for sends	Simpler than WS	Unidirectional; need separate channel for sends

Key Decision: WebSocket for all real-time delivery.
Connection registry in Redis:

```
SET "conn:{user_id}" "chat-server-47" EX 300
```

- Refreshed every 60s by Chat Server via heartbeat.
- If key expires (server crash), user treated as offline → push notification path.

Load balancing WebSocket connections:

- L4 load balancer (not L7) — connections are long-lived TCP; no HTTP termination per request.
- Use consistent hashing on `user_id` for connection routing (optional but reduces cross-server lookups).

6.2 Message Delivery: Online vs Offline

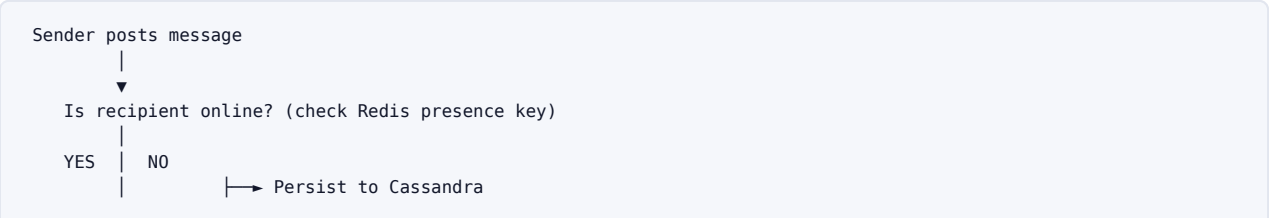

```
        |          └──► Enqueue push notification (APNs/FCM)
        ▼                        |
  Route to recipient's           |
  Chat Server via Redis          ▼
  connection registry    Push Notif Svc sends silent push
        |                → device wakes up, fetches history
        ▼                via REST, renders notification
  Deliver via WebSocket
```

Message deduplication:

- Client generates a UUID (`client_msg_id`) before sending.
- Server stores `client_msg_id` → `server_msg_id` mapping for 24 hours.
- If retry received with same `client_msg_id`, return existing `server_msg_id` (idempotent).

6.3 Group Chat Delivery

Small groups ($\leq$ 500 members):

- Fanout on message write: write one message record, but create one delivery record per member.
- Alternative: write one message, each member's client polls for group updates using a `last_read_seq` cursor.

Recommended for $\leq$ 500 members:

```
Message write → Kafka
Kafka consumer → write 1 message to messages table
               → write N "unread" records to user_conversation table (one per member)
Online members → WebSocket push
Offline members → Push notification
```

Very large groups (> 500, e.g., broadcast channels):

- Switch to fanout-on-read: message is in one place, members' clients pull new messages using `seq_num` cursor.
- Similar to news feed pull model.

6.4 Message Ordering and Sequence Numbers

Problem: Two users send messages simultaneously. Who goes first?

Solution: Sequence number per conversation (not global).

- Each conversation has a monotonic `seq_num` counter stored in Redis or a coordination service (like Zookeeper).
- When Message Service processes a message, it atomically increments `seq_num` for that `conversation_id` and assigns it to the message.
- Messages displayed in `seq_num` order on all clients.

Clock skew: Client timestamps cannot be trusted. Always use server-assigned `seq_num` for ordering. Client timestamp only shown as display time.

```
INCR "seq:{conversation_id}"  → returns 100034 (atomic in Redis)
Store message with seq=100034
```

6.5 Online Presence

Heartbeat protocol:

- Client sends WebSocket heartbeat every 10 seconds.
- Chat Server updates Redis: `SET "presence:{user_id}" 1 EX 30`.
- Key expires in 30s if no heartbeat → user considered offline.

Displaying presence:

- When user A opens a chat with user B: fetch `GET "presence:{user_id_B}"`.
- "Last seen": stored as `last_seen:{user_id}` on heartbeat stop (WebSocket disconnect or timeout).

Scale challenge:

- 500M DAU × heartbeat every 10s = 50M presence writes/sec to Redis.
- **Optimization:** Only send heartbeat if the user is actively viewing a conversation. Idle app = slower heartbeat (30s). Closed app = no heartbeat.
- Use presence fanout sparingly: only push presence updates to contacts who have a conversation open with that user.

6.6 Read Receipts

States: `sent` → `delivered` → `read`.

```
sent:       Message reached the chat server (ack from server to sender).
delivered:  Message delivered to recipient's device (ack from recipient device).
read:       Recipient opened the conversation (client sends "read" ack).
```

Implementation:

- Each message has a `status` field in Cassandra.
- Recipient sends `{"type":"ack","msg_id":"...","ack_type":"read"}` over WebSocket.
- Chat Server → Message Service → update Cassandra status → forward read receipt to sender's Chat Server → sender gets green double-tick.

Group read receipts:

- Store `read_by: Map<user_id, timestamp>` per message in Cassandra.
- "Delivered to all" = all members' devices have acknowledged `delivered`.
- Show "Read by X, Y, Z" as clients send individual read acks.

7. Bottlenecks & Scaling

150 million concurrent WebSocket connections:

- Each Chat Server handles ~50K–100K connections (Go or Node.js event loop).
- Need ~1,500–3,000 Chat Servers.
- Chat Servers are stateful but share no business logic — scale horizontally, no coupling.

Kafka throughput (230K messages/sec):

- Partition by `conversation_id` (ensures ordering per conversation).
- 230K msg/sec × 1 KB = 230 MB/s — Kafka handles this easily with dozens of brokers.
- Message Service consumer group: scale consumers to match Kafka partition count.

Cassandra message storage (3 TB/day):

- Partition by `conversation_id` ensures messages in same conversation are co-located.
- Hot conversation: one partition gets hammered → add `bucket` (time shard) to partition key.
- 3 TB/day × 365 = ~1.1 PB/year. Cassandra compresses ~3:1 → ~370 TB/year of disk.
- TTL old messages: archive messages > 1 year to cold storage (S3 Glacier), keep last 1 year hot.

Media delivery:

- Don't send media over WebSocket — too large.
- Client uploads directly to S3 (pre-signed URL from Media Service).
- Sends media URL in the chat message. Recipients download from CDN-fronted S3.

Push notification scale:

- 500M users. At any time, 70% offline.

- Surge: celebrity group sends a message → 1M push notifications simultaneously.
- Use async workers + APNs/FCM batch APIs. Rate limit per APNs connection (use multiple connections).

8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Real-time transport	**WebSocket**	Long polling / SSE	Full-duplex; lowest latency; widely supported
Message storage	**Cassandra (conv-partitioned)**	MySQL (sharded)	Write throughput; time-series access pattern; no joins needed
Connection registry	**Redis (TTL-based)**	Consistent hash routing	Fast lookup; auto-evicts stale connections on TTL expiry
Message ordering	**Server-assigned seq_num (Redis INCR)**	Vector clocks	Simple; sufficient for single-conversation ordering
Group delivery	**Fanout to members (< 500) / pull (> 500)**	Always pull	Balances latency vs write amplification by group size
Presence heartbeat	**10s client heartbeat, 30s TTL**	Server-push disconnect events	Handles silent disconnects (network failure, app killed)
Read receipts	**Per-message status in Cassandra**	Separate receipt table	Avoids join; co-located with message data
Media	**S3 + CDN (URL in message)**	In-message blob	Media over WebSocket is impractical at scale

End of Part 1 — Foundational System Design Problems.