

System Design Problems — Foundational



In This Edition

1	1. Design a URL Shortener (TinyURL / bit.ly)	2
1.1	1. Requirements	2
1.2	2. Capacity Estimation	2
1.3	3. API Design	3
1.4	4. Data Model	3
1.5	5. High-Level Architecture	4
1.6	6. Deep Dives	5
1.7	7. Bottlenecks & Scaling	6
1.8	8. Trade-offs Summary	7
2	2. Design a Rate Limiter	7
2.1	1. Requirements	8
2.2	2. Capacity Estimation	8
2.3	3. API Design	9
2.4	4. Data Model	10
2.5	5. High-Level Architecture	10
2.6	6. Deep Dives	11
2.7	7. Bottlenecks & Scaling	13
2.8	8. Trade-offs Summary	13
3	3. Design a News Feed (Facebook/Twitter Timeline)	13
3.1	1. Requirements	14
3.2	2. Capacity Estimation	14
3.3	3. API Design	15
3.4	4. Data Model	15
3.5	5. High-Level Architecture	16
3.6	6. Deep Dives	17
3.7	7. Bottlenecks & Scaling	18
3.8	8. Trade-offs Summary	18
4	4. Design a Chat System (WhatsApp / Messenger)	19
4.1	1. Requirements	19
4.2	2. Capacity Estimation	19
4.3	3. API Design	20
4.4	4. Data Model	20
4.5	5. High-Level Architecture	21
4.6	6. Deep Dives	22
4.7	7. Bottlenecks & Scaling	25
4.8	8. Trade-offs Summary	25

Complete end-to-end worked designs for a 45–60 min FAANG system-design round. Each section mirrors the structure you should drive in the actual interview.

1 1. Design a URL Shortener (TinyURL / bit.ly)

1.1 1. Requirements

Clarifying questions to ask the interviewer:

- › What is the expected scale? (DAU, writes per second, reads per second?)
- › Should we support custom aliases (user-defined short codes)?
- › What is the required URL lifetime? Do links expire?
- › Do we need analytics (click counts, geo, referrer)?
- › Should we support link editing or deletion?
- › Is high availability or strong consistency more important?

Functional Requirements:

1. Given a long URL, return a unique short URL (e.g., `https://tny.io/aB3kX9`).
2. Redirect a short URL to the original long URL.
3. (Optional) Custom aliases: user can specify the short code.
4. (Optional) URL expiry: links expire after N days.
5. (Optional) Analytics: track click count, geography, referrer.

Non-Functional Requirements:

- › 100 ms P99 redirect latency.
- › 99.99% availability (URL shortener is a critical dependency for customers).
- › Short codes must be unique and not predictable/enumerable.
- › The system is **read-heavy** (100:1 read-to-write ratio is typical).

1.2 2. Capacity Estimation

Assumptions:

- › 500 million URLs shortened per month (writes).
- › 100:1 read-to-write ratio → 50 billion redirects per month.

Write QPS:

$500,000,000 / (30 \times 86,400) \approx 193 \text{ writes/sec} \approx \sim 200 \text{ writes/sec}$
 Peak (3x): $\approx \sim 600 \text{ writes/sec}$

Read QPS:

$50,000,000,000 / (30 \times 86,400) \approx 19,290 \text{ reads/sec} \approx \sim 20,000 \text{ reads/sec}$
 Peak (3x): $\approx \sim 60,000 \text{ reads/sec}$

Storage (URLs, 5 years):

```

1 500M writes/month × 12 months × 5 years = 30 billion URLs
2 Each record ≈ 500 bytes (Long URL ~400 B + metadata ~100 B)
3 Total: 30B × 500 B = 15 TB
  
```

Cache (top 20% URLs serve 80% traffic):

$20,000 \text{ reads/sec} \times 86,400 \text{ sec/day} = 1.728 \text{ billion reads/day}$
 20% of 30B URLs = 6 billion "hot" URLs $\rightarrow$ not all fit in RAM
 Pragmatic: cache last 24 hrs of writes = $500\text{M}/30 \times 500\text{B} \approx 8.3 \text{ GB/day} \rightarrow \sim 10 \text{ GB RAM}$

Bandwidth:

Inbound (writes): $200/\text{sec} \times 500 \text{ B} \approx 100 \text{ KB/s}$
 Outbound (reads): $20,000/\text{sec} \times 500 \text{ B} \approx 10 \text{ MB/s}$

1.3 3. API Design

```

1 # Shorten a URL
2 POST /api/v1/shorten
3 Request:
4 {
5     "long_url": "https://www.example.com/very/long/path?q=something",
6     "custom_code": "my-alias", // optional
7     "expires_at": "2027-01-01T00:00Z" // optional
8 }
9 Response 201:
10 {
11     "short_url": "https://tny.io/aB3kX9",
12     "short_code": "aB3kX9",
13     "expires_at": "2027-01-01T00:00Z"
14 }
15
16 # Redirect (the critical read path)
17 GET /{short_code}
18 Response 302:
19 Location: https://www.example.com/very/long/path?q=something
20
21 # Analytics (optional)
22 GET /api/v1/stats/{short_code}
23 Response 200:
24 {
25     "clicks": 14982,
26     "top_countries": [...],
27     "referrers": [...]
28 }
29
30 # Delete
31 DELETE /api/v1/shorten/{short_code}
32 Response 204
  
```

1.4 4. Data Model**SQL vs NoSQL decision:**

- › The primary access pattern is a simple key-value lookup: `short_code` $\rightarrow$ `long_url`.
- › No complex joins needed; writes are straightforward inserts.
- › **Choice: NoSQL key-value store (Cassandra or DynamoDB) for the URL table** — horizontal scale, high availability, and the access pattern is a perfect fit.
- › A relational DB (PostgreSQL) is used for user/account data where consistency matters.

URL Table (Cassandra / DynamoDB):

```

1 urls
2
3 short_code    TEXT        PRIMARY KEY    -- "aB3kX9"
4 long_url     TEXT        NOT NULL       -- original URL
5 user_id      UUID
6 created_at   TIMESTAMP
7 expires_at   TIMESTAMP          -- NULL = never expires
8 is_deleted   BOOLEAN    DEFAULT false

```

Users Table (PostgreSQL):

```

users

user_id      UUID        PRIMARY KEY
email        VARCHAR(255) UNIQUE NOT NULL
api_key      VARCHAR(64) UNIQUE          -- for rate limiting
created_at   TIMESTAMP

```


Analytics Table (write-optimized, e.g., Cassandra or ClickHouse):


```

1 click_events
2
3 short_code    TEXT
4 clicked_at    TIMESTAMP
5 country       TEXT
6 referrer      TEXT
7 user_agent    TEXT
8 -- Partition key: short_code; Clustering key: clicked_at DESC

```

1.5 5. High-Level Architecture





Component walkthrough:

1. **Client** sends POST /shorten OR GET /{code} over HTTPS.
2. **Load Balancer** routes to stateless Write Service or Redirect Service pods.
3. **Redirect Service** checks **Redis** first (cache hit → 302 immediately). On cache miss, reads Cassandra, populates cache, then redirects.
4. **Write Service** calls **KGS** (Key Generation Service) for a pre-generated unique code, stores (code, long_url) in Cassandra, and warms the cache.
5. **Kafka** asynchronously receives click events; **Analytics Service** batch-writes to ClickHouse for reporting.

1.6 6. Deep Dives

6.1 Key Generation: How to create short codes?

Approach	How	Pros	Cons
MD5 / SHA-256 hash	Hash long URL, take first 6–8 chars	Deterministic (same URL → same code)	Collisions possible; must check DB on every write
Counter + Base62	Global atomic counter → encode in Base62	No collisions; simple	Single counter = SPOF; reveals ordering/volume
Pre-generated KGS	Dedicated service pre-generates & stores unused keys in a "keys" table	No collision checking at write time; fast	Additional service to maintain; key theft if service crashes
UUID	Generate random UUID, take first 8 chars	Simple	High collision probability at scale

Key Decision: KGS (Key Generation Service) is the recommended approach for large scale.

KGS design:

- KGS pre-generates 6-character Base62 codes ($62^6 \approx 56$ billion unique codes).
- Stores them in two tables: `unused_keys` and `used_keys`.
- Each Write Service instance requests a batch (e.g., 1,000 keys) from KGS on startup, holds them in-memory.
- KGS marks those keys as "in use" atomically. If a Write Service pod crashes, those keys are wasted but never reused — acceptable loss.
- KGS itself can be replicated with a leader for key assignment.

6.2 Redirect: 301 vs 302?

	301 Permanent	302 Temporary
Browser caches?	Yes — browser won't call TinyURL again	No — browser calls TinyURL on every click
Analytics accuracy	Broken — browser bypasses your servers	Accurate — every click hits your servers
Load on TinyURL	Lower after first visit	Higher (every click)
Recommendation	Use only if analytics not needed	Use 302 for full analytics capability

Interview takeaway: Always use 302 unless you explicitly trade analytics for reduced server load.

6.3 Collision Handling

If using hash-based approach:

1. Compute MD5(long_url) → take first 7 chars as short_code.
2. Check if short_code exists in DB.
3. If collision: append a salt (e.g., increment integer), re-hash, retry.
4. Limit retries to 3; fall back to random generation if exhausted.

With KGS, collisions are impossible — keys are pre-assigned and marked used atomically.

6.4 Custom Aliases

- › User provides desired short code (e.g., my-brand).
- › Check availability: `SELECT 1 FROM urls WHERE short_code = 'my-brand'.`
- › If taken → return 409 Conflict.
- › Store with `is_custom = true` flag to distinguish from generated codes.
- › Rate-limit custom alias creation per user to prevent squatting.

6.5 URL Expiry

- › Store `expires_at` timestamp in the URL record.
- › On redirect: if `expires_at < NOW()` → return 404 (or 410 Gone).
- › Background cleanup job: delete expired records weekly (keep DB lean).
- › Cache entries: set Redis TTL = `expires_at - NOW()` so expired URLs auto-evict.

1.7 7. Bottlenecks & Scaling

Read hotspot — viral URLs:

- › A single URL (e.g., link from a viral tweet) can receive millions of reads/sec.
- › Redis handles this well — single key lookup is $O(1)$ and Redis handles 100k+ ops/sec per node.
- › For extreme cases: replicate the specific key across multiple Redis nodes or use local in-process cache (LRU with 1000 entries) in each Redirect Service pod.

Write scaling:

- › Writes at ~200/sec are trivially handled. Cassandra scales horizontally; add nodes as needed.
- › KGS: single KGS with a read replica for HA is sufficient at this scale. Key generation is CPU-light.

DB sharding:

- › Cassandra partitions by short_code automatically (consistent hashing). No manual sharding needed.

- › If using SQL: shard by `short_code` modulo `N` (hash partitioning). Use a consistent hash ring to avoid resharding pain.

Cache eviction:

- › Redis with LRU policy. Set max memory = 10–20 GB.
- › TTL on each entry = `min(expires_at, 24 hrs)` to prevent stale data.

Failure handling:

- › Redis down → fall through to Cassandra (latency increases but system stays up).
- › KGS down → Write Service uses its in-memory key batch until KGS recovers; alert if batch exhausted.
- › Cassandra node down → Cassandra replication factor 3 (RF=3) ensures reads/writes still succeed with 2 nodes up.

Geographic distribution:

- › Deploy read replicas of Cassandra in multiple regions.
- › Use GeoDNS to route users to nearest Redirect Service cluster.
- › URL writes go to a primary region; replicated asynchronously to others.

1.8 8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Key generation	KGS (pre-generated)	Hash + collision check	No collision risk; predictable latency
DB for URLs	Cassandra (NoSQL)	PostgreSQL	Write scalability; simple key-value access pattern
Redirect type	302 Temporary	301 Permanent	Enables accurate analytics
Cache	Redis (LRU)	Memcached	Rich data types; TTL support; persistence
Analytics pipeline	Kafka → ClickHouse (async)	Sync write to SQL	Decouples write path; handles bursty click events
Custom aliases	Supported (rate-limited)	Not supported	Product differentiator; must prevent abuse

2 2. Design a Rate Limiter




Figure 1. A bucket refills at a steady rate up to a capacity. Each request spends a token; an empty bucket throttles. This smooths traffic while allowing short bursts.

2.1 1. Requirements

Clarifying questions to ask the interviewer:

- › Where should the rate limiter live? Client-side, API gateway, or per-service middleware?
- › What are we rate-limiting on? Per user? Per IP? Per API key? Per endpoint?
- › What happens when a limit is exceeded — hard reject (429) or soft throttle (queue)?
- › Do we need distributed rate limiting (multiple servers sharing state)?
- › What are the latency requirements? Rate limit check must not add significant overhead.
- › Do rules need to be configurable at runtime without deployments?

Functional Requirements:

1. Limit the number of requests a client (user/IP/API key) can make in a time window.
2. Return HTTP 429 Too Many Requests when limit is exceeded, with Retry-After header.
3. Support multiple limits (e.g., 100 req/min AND 1,000 req/hr for same client).
4. Rules are configurable without service restarts.

Non-Functional Requirements:

- › Rate limit check must add < 5 ms P99 latency overhead.
- › Highly available — rate limiter failure should fail open (allow traffic) not fail closed (block everything).
- › Works across multiple servers (distributed, consistent limits).
- › Accurate enough — exact precision not required; small over-limit acceptable.

2.2 2. Capacity Estimation

Assumptions:

- › 10 million DAU, each making ~100 requests/day → 1 billion requests/day.
- › Rate limit state must be checked for every request.

QPS to rate limiter:

$1,000,000,000 / 86,400 \approx 11,574 \text{ req/sec} \approx \sim 12,000 \text{ checks/sec}$
 Peak (3×): $\approx \sim 36,000 \text{ checks/sec}$

Memory for rate limit counters:

Strategy: sliding window counter per user per endpoint, 1-hr window.
 $10M \text{ users} \times 10 \text{ endpoints} \times (\text{counter} + \text{timestamp}) \approx 10M \times 10 \times 20 \text{ bytes} = 2 \text{ GB}$
 → Fits comfortably in a Redis cluster.

Bandwidth:

- › Rate limit check: small — just a counter lookup and increment.
- › Each Redis call $\approx$ microseconds; no large data transfer.

2.3 3. API Design

The rate limiter is typically middleware, not a standalone API. But if exposing as a service:

```

1 # Check and record a request (called by API gateway)
2 POST /rate-limit/check
3 Request:
4   {
5     "client_id": "user:1234",
6     "endpoint": "/api/v1/posts",
7     "timestamp": 1718000000
8   }
9 Response 200 (allowed):
10  {
11    "allowed": true,
12    "remaining": 42,
13    "reset_at": 1718003600
14  }
15 Response 429 (throttled):
16  {
17    "allowed": false,
18    "retry_after": 58,           // seconds until window resets
19    "limit": 100,
20    "window": "1m"
21  }
22
23 # Update rules (admin)
24 PUT /rate-limit/rules/{client_id}
25 Request:
26   {
27     "rules": [
28       { "window": "1m", "limit": 100 },
29       { "window": "1h", "limit": 1000 }
30     ]
31   }
  
```

Response headers to include on every API response:

X-RateLimit-Limit: 100
 X-RateLimit-Remaining: 42
 X-RateLimit-Reset: 1718003600
 Retry-After: 58 // only on 429

2.4 4. Data Model

Redis data structures (no SQL needed — state is ephemeral):
Token Bucket / Fixed Window counter:

```
1 Key: "rl:{client_id}:{endpoint}:{window_start}"
2 Value: integer (request count)
3 TTL: window duration (e.g., 60s for 1-min window)
```


Sliding Window Log:

```
1 Key: "rl:{client_id}:{endpoint}"
2 Type: Redis Sorted Set
3 Member: request_id (UUID)
4 Score: timestamp (epoch ms)
5 TTL: window duration
```


Rules config (read rarely, cache in-process):

```
1 # Redis Hash or a config store (Consul / etcd)
2 Key: "rule:{client_id}"
3 Value: JSON { "rules": [{"window": "1m", "limit": 100}, ...] }
```

2.5 5. High-Level Architecture



Rules flow:

**Component walkthrough:**

1. Every request hits the **API Gateway**.
2. Gateway invokes **Rate Limiter Middleware** (same process or sidecar).
3. Middleware reads the rule for (client_id, endpoint) from its **in-process cache** (refreshed every 30s from etcd/Consul).
4. Middleware atomically increments/checks the counter in **Redis**.
5. If allowed: forward request to **Backend Service**.
6. If denied: return 429 immediately with Retry-After.

2.6 6. Deep Dives

6.1 Algorithm Comparison

Algorithm	How it works	Pros	Cons	Best for
Fixed Window Counter	Reset counter every N seconds	Simple; O(1) memory	Boundary burst: $2 \times$ limit possible at window edge	Simple internal limits
Sliding Window Log	Store timestamp of every request in sorted set; count within window	Accurate	O(n) memory per user (n = requests in window)	Low-volume, high-accuracy needs
Sliding Window Counter	Blend current + previous window counts weighted by overlap	Near-accurate, O(1) memory	Slight inaccuracy (~0.003% over-limit)	Best general choice
Token Bucket	Bucket of tokens refilled at rate R; request consumes one	Allows bursts up to bucket size	Harder to reason about exact limits; burst may confuse users	APIs that allow bursting
Leaky Bucket	Queue requests; process at constant rate	Smooth output rate	Queue introduces latency; bursty requests get delayed	Network traffic shaping

Key Decision: Sliding Window Counter for most API rate limiting. Token Bucket when you want to allow short bursts.

6.2 Sliding Window Counter --- Implementation Detail

```

Window = 1 minute.
Current minute bucket: starts at :00
Previous minute bucket: started at :-60

count = prev_bucket_count * (overlap / window) + curr_bucket_count

Example at :45 into the current minute:
overlap = 60 - 45 = 15 seconds (15/60 = 25% of previous window still counts)
count = prev_count * 0.25 + curr_count

If count > limit → reject.
  
```

Redis implementation:

```

1  -- Lua script (atomic)
2  local key_curr = KEYS[1] -- "rl:user1:1718000060"
3  local key_prev = KEYS[2] -- "rl:user1:1718000000"
4  local limit   = tonumber(ARGV[1])
5  local now     = tonumber(ARGV[2])
6  local window  = tonumber(ARGV[3])
7  local window_start = now - (now % window)
8  local overlap  = window - (now - window_start)
9
10 local curr_count = tonumber(redis.call('GET', key_curr) or 0)
11 local prev_count = tonumber(redis.call('GET', key_prev) or 0)
12 local weighted  = prev_count * (overlap / window) + curr_count
13
14 if weighted >= limit then
15     return 0 -- rejected
16 end
17 redis.call('INCR', key_curr)
18 redis.call('EXPIRE', key_curr, window * 2)
19 return 1 -- allowed

```

6.3 Where to Place the Rate Limiter?

Placement	Pros	Cons
Client-side	Zero server load	Easily bypassed; untrustworthy
API Gateway	Single enforcement point; language-agnostic	Gateway becomes a bottleneck; less context
Service middleware	Per-service granularity; business context	Duplicated logic across services; harder to change rules centrally
Dedicated Rate Limit Service	Centralized rules; all services call it	Extra network hop; additional failure point
Recommended: Gateway + Redis	Rules centralized in Redis; enforced at gateway	Need to handle Redis failure gracefully

6.4 Distributed Rate Limiting --- Race Conditions

Problem: Two requests arrive simultaneously on different servers. Both read count=99 (limit=100), both increment to 100, both allow. Result: 2 requests slip through over limit.

Solutions:

1. **Redis Lua scripts (recommended):** Atomic read-increment-check in a single Lua script. Redis is single-threaded; no race condition possible.
2. **Redis INCR + EXPIRE pattern:**

```

1  count = INCR(key)
2  if count == 1: EXPIRE(key, window) -- set TTL only on first request
3  if count > limit: return 429

```

Problem: race between INCR and EXPIRE if server crashes. Mitigate with SET key 0 EX window NX on first call.

3. **Sticky sessions / consistent hashing:** Route same client_id to same rate limiter node. Eliminates distributed counting but adds routing complexity and uneven load.

Key Decision: Use Redis Lua scripts for atomicity. No locks, no transactions overhead.

6.5 Fail-Open vs Fail-Closed

- › **Fail-open (recommended):** If Redis is down, allow all requests. Availability > perfect rate limiting.
- › **Fail-closed:** If Redis is down, reject all requests. Too severe for most use cases — takes down your API.
- › Implementation: wrap Redis call in try-catch; on exception, log and allow the request, alert on-call.

2.7 7. Bottlenecks & Scaling

Redis as single point of failure:

- › Use Redis Cluster (sharded) — each client_id maps to a shard via consistent hashing.
- › Redis Sentinel for HA (automatic failover).
- › Replication lag between primary and replica: acceptable (small over-limit possible).

Redis latency:

- › Redis P99 $\approx$ 0.5–1 ms in-datacenter. Total rate limit overhead < 2 ms.
- › If Redis is co-located on same host (Redis sidecar per gateway node): P99 < 0.2 ms.

Rule changes at scale:

- › Store rules in etcd/Consul. Rate limiter middleware subscribes to watch-notifications.
- › In-process cache with 30s TTL: small window where old rules apply after update.

Hotspots (one user generating massive traffic):

- › Rate limiter itself handles this well — a single Redis key gets hammered, but Redis is designed for this.
- › Use local in-process counter as first layer (no Redis call for clearly non-bursting clients), then Redis for enforcement.

2.8 8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Algorithm	Sliding Window Counter	Token Bucket	Near-exact accuracy; O(1) memory; prevents boundary bursts
State store	Redis Cluster	In-process memory	Distributed consistency across servers; fast enough
Placement	API Gateway middleware	Dedicated service	Avoids extra network hop; centralizes enforcement
Atomicity	Redis Lua scripts	Redis WATCH/MULTI	Simpler; no retry logic needed
Failure mode	Fail-open	Fail-closed	Availability > perfect accuracy for most systems
Rules storage	etcd + in-process cache	DB query per request	Rules rarely change; in-process cache eliminates lookup overhead

3 3. Design a News Feed (Facebook/Twitter Timeline)

3.1 1. Requirements

Clarifying questions to ask the interviewer:

- › Is this a "friends" model (Facebook — bidirectional) or "followers" model (Twitter — unidirectional)?
- › What is the feed ordering — chronological or ranked by relevance?
- › What's the expected social graph size? Max followers per user?
- › Do we support stories (ephemeral) vs. posts (permanent)?
- › Is the feed pre-generated (push) or assembled on request (pull)?
- › Do we need to support ads injection in the feed?

Functional Requirements:

1. User can create a post (text, image, video link).
2. User sees a feed of posts from people they follow, ordered by recency or ranking.
3. Feed updates in near-real-time (within seconds of post creation).
4. Pagination: load N posts, scroll for more.
5. Users can like, comment, share (out of scope for this design; focus on feed generation).

Non-Functional Requirements:

- › 500 million DAU (Facebook scale).
- › Feed load latency < 200 ms.
- › Post creation propagates to followers within 5 seconds (eventual consistency is OK).
- › Users can have up to 5,000 friends (Facebook) or millions of followers (Twitter celebrities).

3.2 2. Capacity Estimation

Assumptions:

- › 500 million DAU.
- › Each user views feed 5 times/day → 2.5 billion feed reads/day.
- › Each user creates 1 post every 5 days → 100 million new posts/day.
- › Average friends/followers: 300.

Post write QPS:

$100,000,000 / 86,400 \approx 1,157 \text{ writes/sec} \approx \sim 1,200 \text{ posts/sec}$
 Peak (5x): $\approx \sim 6,000 \text{ posts/sec}$

Feed read QPS:

$2,500,000,000 / 86,400 \approx 28,935 \text{ reads/sec} \approx \sim 29,000 \text{ reads/sec}$
 Peak (3x): $\approx \sim 87,000 \text{ reads/sec}$

Fanout writes (push model — writing to all follower feeds on post):

$1,200 \text{ posts/sec} \times 300 \text{ avg followers} = 360,000 \text{ feed writes/sec}$
 Celebrity (10M followers): 1 post → 10M writes (burst)

Storage:

Posts: $100\text{M/day} \times 1 \text{ KB (text)} = 100 \text{ GB/day text}$
 Images/videos: separate blob store (S3)
 Feed cache: $500\text{M users} \times 20 \text{ cached posts} \times 1 \text{ KB} = 10 \text{ TB} \rightarrow \text{too big for RAM}$
 → Cache only active users (10% daily active at peak):

50M users × 20 posts × 1 KB = 1 TB → feasible with Redis cluster

3.3 3. API Design

```
1 # Create a post
2 POST /api/v1/posts
3 Auth: Bearer token
4 Request:
5 {
6     "content": "Hello world!",
7     "media_urls": ["https://cdn.example.com/img/abc.jpg"],
8     "visibility": "friends" // or "public", "private"
9 }
10 Response 201:
11 {
12     "post_id": "p_9x3kA2",
13     "created_at": "2026-06-14T10:00:00Z"
14 }
15
16 # Get news feed
17 GET /api/v1/feed?limit=20&cursor=<opaque_cursor>
18 Auth: Bearer token
19 Response 200:
20 {
21     "posts": [
22         {
23             "post_id": "p_9x3kA2",
24             "author_id": "u_4f2k",
25             "author_name": "Alice",
26             "content": "Hello world!",
27             "media_urls": [...],
28             "like_count": 142,
29             "comment_count": 7,
30             "created_at": "2026-06-14T10:00:00Z"
31         },
32         ...
33     ],
34     "next_cursor": "<opaque_cursor>"
35 }
36
37 # Get a user's profile posts
38 GET /api/v1/users/{user_id}/posts?limit=20&cursor=<cursor>
39 Response 200: { "posts": [...], "next_cursor": "..." }
```

3.4 4. Data Model

SQL vs NoSQL:

- **Posts table:** Write-once, read-many. High read QPS. Shard by `author_id`. → NoSQL (Cassandra) or sharded PostgreSQL.
- **Social graph (follows):** Relationship lookups ("who does user X follow?"). → Graph DB (Neo4j) or sharded SQL; at Facebook scale: custom TAO (graph store).
- **Feed cache:** Redis sorted sets (score = post timestamp).
- **User data:** Relational (PostgreSQL) — low write volume, complex queries.

Posts Table (Cassandra — partition by author, cluster by time):

posts

```
author_id  UUID      PARTITION KEY
post_id    TIMEUUID  CLUSTERING KEY (DESC)  -- newest first
content    TEXT
media_urls LIST<TEXT>
like_count COUNTER
created_at TIMESTAMP
```

Social Graph Table (Cassandra — two tables for O(1) lookups):

follows_by_follower -- "who does user X follow?"

```
follower_id  UUID  PARTITION KEY
followee_id  UUID  CLUSTERING KEY
```


follows_by_followee -- "who follows user X?"

```
followee_id  UUID  PARTITION KEY
follower_id  UUID  CLUSTERING KEY
```

Feed Cache (Redis Sorted Set per user):

```
1 Key:      "feed:{user_id}"
2 Members: post_id
3 Scores:   timestamp (epoch ms) — higher = newer
4 TTL:      48 hours (evict inactive users)
5 Max size: 800 posts per user (trim older entries)
```

3.5 5. High-Level Architecture



Post creation flow: Post Svc → Cassandra + Kafka
 Fanout flow: Kafka → Fanout Workers → Redis (write to follower feeds)
 Feed read flow: Feed Svc → Redis cache → (miss) Cassandra

3.6 6. Deep Dives

6.1 Fanout Strategies

Strategy	How	Pros	Cons	Best for
Fanout-on-Write (Push)	On post creation, write post_id to each follower's feed cache	Feed reads O(1) — just read from cache	Write amplification: 1 post → N writes; celebrity (10M followers) = 10M cache writes	Average users
Fanout-on-Read (Pull)	On feed request, fetch posts from each followee, merge, sort	No write amplification	Feed read O(N followees) — slow for users following many people	Celebrity accounts, inactive followers
Hybrid (recommended)	Push to regular followers; pull for celebrities at read time; don't push to inactive users	Balances read and write load	More complex logic	Large-scale production systems

Key Decision: Hybrid fanout.

Implementation:

- › Define "celebrity" as users with > 1M followers.
- › On post write: fanout-on-write for non-celebrities; for celebrities, mark the post in a "celebrity posts" index only.
- › On feed read: read user's pre-built feed from Redis (push-mode posts) + fetch latest celebrity posts from celebrity index + merge and deduplicate.
- › Skip fanout for users who haven't been active in 7 days (check last_active timestamp).

6.2 Celebrity Problem (Hotspot)

A celebrity with 10 million followers posts → need 10 million cache writes within seconds.

Mitigation:

1. **Async fanout workers:** Put the fanout job on Kafka with 10,000 partitions. Workers consume in parallel — 10M writes spread across workers.
2. **Pull for celebrity content:** Don't write celebrity posts to follower feeds. At read time, the top-K celebrities a user follows are fetched separately and injected into the feed.
3. **Pre-fetch:** For celebrities, replicate their latest posts to a separate, highly-cached "celebrity feed" shard. All followers read from the same shard (CDN-like effect).

6.3 Feed Ranking

Simple timeline: sort by score = timestamp.

ML-ranked feed (Facebook/Instagram model):

```

1 score = f(recency, engagement_velocity, relationship_strength, content_type)
2
3 engagement_velocity = likes_per_hour / avg_likes_per_hour_for_author
4 relationship_strength = interaction_history(viewer, author)
```

Interview takeaway: In the interview, acknowledge ranking but defer details. Say: "We'd train an ML model offline, serve predictions via a Feature Store + low-latency inference service. The feed service re-ranks the top 200 candidates before returning the top 20."

6.4 Pagination

Cursor-based pagination (not offset):

- Offset pagination (`LIMIT 20 OFFSET 100`) breaks when new posts are inserted — items shift, causing duplicates/skips.
- Cursor = encoded last-seen post timestamp + post_id.
- Redis: `ZRANGEBYSCORE feed:{user_id} -inf <cursor_score> LIMIT 0 20`.
- Cursor makes the feed stable even as new posts arrive.

3.7 7. Bottlenecks & Scaling

Feed cache size:

- 10 TB for active users — run 50+ Redis nodes (200 GB each) in a cluster.
- Use consistent hashing for user → Redis shard mapping.
- Evict feeds of users inactive for 7 days to save memory.

Cassandra post storage:

- Partition by author_id → hot partition if one author posts very frequently.
- Mitigate: add bucket to partition key: (author_id, bucket) where bucket = created_at / 30 days.

Social graph at scale:

- Facebook built TAO (The Associations and Objects) — a custom graph store on top of MySQL + Memcache.
- At smaller scale: Cassandra follows_by_follower + follows_by_followee tables work well.

Read path latency breakdown (target < 200 ms):

```
Redis feed read:           5 ms
Celebrity post fetch:     20 ms (parallel)
User profile hydration:   10 ms (from user cache)
Ranking re-score:        15 ms
Network + serialization: 10 ms
Total:                   ~60 ms (well within 200 ms budget)
```

3.8 8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Fanout strategy	Hybrid (push+pull)	Pure push or pure pull	Balances write amplification vs read latency
Feed storage	Redis sorted sets	DB query at read time	Sub-10ms reads; DB can't handle 87K reads/sec directly
Feed ordering	Reverse-chronological (MVP)	ML-ranked feed	Simpler; ranked feed is a layer on top
Pagination	Cursor-based	Offset-based	Stable under inserts; no duplicate/skip on new posts
Post DB	Cassandra	PostgreSQL (sharded)	Write throughput; natural time-series clustering

Decision	Choice Made	Alternative	Reason
Celebrity handling	Pull at read time	Fanout to all followers	Avoids 10M+ write storms per celebrity post

4 4. Design a Chat System (WhatsApp / Messenger)

4.1 1. Requirements

Clarifying questions to ask the interviewer:

- › 1:1 chat only, or also group chat? What is the max group size?
- › Do we need real-time delivery (online) and push notifications (offline)?
- › Media support (images, video, voice)?
- › Do messages need end-to-end encryption?
- › Message history: how far back? Searchable?
- › Read receipts (sent / delivered / read)?
- › Online presence ("last seen")?

Functional Requirements:

- 1:1 messaging with real-time delivery.
- Group messaging (up to 500 members per group).
- Online/offline status with "last seen" timestamp.
- Read receipts: sent → delivered → read.
- Push notifications for offline users.
- Message history persisted; searchable.
- Media: images, files (< 100 MB each).

Non-Functional Requirements:

- › 500 million DAU; 100 billion messages/day.
- › Message delivery latency < 500 ms (online users).
- › High availability: 99.99%.
- › Messages delivered exactly once and in order.
- › End-to-end encryption support (Signal Protocol) — note architecture impact.

4.2 2. Capacity Estimation

Assumptions:

- › 500 million DAU.
- › Each user sends 40 messages/day → 20 billion messages/day (1:1 + group).
- › Average message size: 100 bytes (text).
- › Average group has 10 members → group message creates 9 additional deliveries.
- › Media: 5% of messages are media → 1 billion media messages/day.

Message write QPS:

20,000,000,000 / 86,400 ≈ 231,481 msg/sec ≈ ~230,000 messages/sec
 Peak (3×): ≈ ~700,000 messages/sec

Storage:

```

Text messages: 20B/day × 100 B      = 2 TB/day
Media:         1B/day × 1 MB avg    = 1 PB/day → S3/GCS
Metadata:      20B/day × 50 B      = 1 TB/day
Total text+meta ≈ 3 TB/day → 1.1 PB/year (text only)

```

Concurrent WebSocket connections:

```

500M DAU × 30% online simultaneously = 150M concurrent connections
Each connection: 65KB OS overhead
Total: 150M × 65KB ≈ 9.75 TB RAM just for connections
→ Need ~100,000 chat servers @ 1,500 connections each
(In practice: use event-driven servers; Nginx/Go handles 100K connections per server)
→ 1,500 servers (100K connections each)

```

4.3 3. API Design

```

# WebSocket connection (real-time channel)
ws://chat.example.com/ws?token=<auth_token>

# WebSocket message format (client → server)
{
  "type": "message",
  "msg_id": "client-generated-uuid", // for deduplication
  "to": "user:5678", // or "group:9999"
  "content": "Hey!",
  "media_url": null,
  "timestamp": 1718000000000 // client time ms
}

# WebSocket message format (server → client, push delivery)
{
  "type": "message",
  "msg_id": "server-assigned-uuid",
  "from": "user:1234",
  "to": "user:5678",
  "content": "Hey!",
  "server_ts": 1718000000050, // server timestamp (canonical ordering)
  "seq": 100034 // sequence number in conversation
}

# ACK (client → server, delivery acknowledgement)
{ "type": "ack", "msg_id": "...", "ack_type": "delivered" }
{ "type": "ack", "msg_id": "...", "ack_type": "read" }

# REST: Send message (fallback for HTTP clients)
POST /api/v1/messages
Request: { "to": "user:5678", "content": "Hey!" }
Response: { "msg_id": "...", "server_ts": 1718000000050 }

# REST: Fetch message history
GET /api/v1/conversations/{conversation_id}/messages?before=<msg_id>&limit=50
Response: { "messages": [...], "has_more": true }

# REST: Upload media
POST /api/v1/media/upload
Response: { "media_url": "https://cdn.example.com/m/abc.jpg", "expires_in": 86400 }

# REST: Get online status
GET /api/v1/users/{user_id}/presence
Response: { "online": false, "last_seen": "2026-06-14T09:55:00Z" }

```

4.4 4. Data Model

Storage choices:

- › **Messages:** HBase / Cassandra — write-heavy time-series; partition by conversation, cluster by time.

- › **Conversations/Groups:** Cassandra or DynamoDB.
- › **User data:** PostgreSQL (low write, complex queries).
- › **Presence (online status):** Redis (ephemeral, TTL-based).
- › **Message queue for offline delivery:** Kafka or internal queues.
- › **Media:** S3 / GCS (blob).

Messages Table (HBase / Cassandra):

1	messages		
2			
3	conversation_id	TEXT	PARTITION KEY -- "conv:{user1}_{user2}" or "group:{id}"
4	seq_num	BIGINT	CLUSTERING KEY DESC -- monotonically increasing per conversation
5	msg_id	UUID	-- globally unique
6	sender_id	UUID	
7	content	TEXT	
8	media_url	TEXT	
9	msg_type	TEXT	-- "text", "image", "video", "system"
10	status	TEXT	-- "sent", "delivered", "read"
11	created_at	TIMESTAMP	

Conversations Table:

conversations		
conversation_id	TEXT	PRIMARY KEY
participants	LIST<UUID>	
group_name	TEXT	(null for 1:1)
last_msg_id	UUID	
last_msg_at	TIMESTAMP	


User Conversations Index (reverse lookup):


user_conversations	-- "what conversations is user X in?"	
user_id	UUID	PARTITION KEY
conversation_id	TEXT	CLUSTERING KEY
last_msg_at	TIMESTAMP	(for sorting)
unread_count	INT	

Presence Table (Redis):

1	Key: "presence:{user_id}"
2	Value: { "online": true, "last_heartbeat": 1718000000000 }
3	TTL: 30 seconds (refreshed by heartbeat every 10s)

4.5 5. High-Level Architecture





Message delivery flow (1:1, both online):

1. Sender client sends message over WebSocket to **Chat Server A**.
2. Chat Server A writes message to **Kafka** (topic: messages).
3. **Message Service** consumes from Kafka, persists to Cassandra.
4. Message Service looks up which **Chat Server** the recipient is connected to (via Redis connection registry: "ws_server:{user_id}" → "chat-server-B").
5. Message Service pushes message to **Chat Server B** (via internal pub/sub).
6. Chat Server B delivers to recipient's WebSocket.
7. Recipient client sends ack: delivered → Chat Server B → Message Service → Cassandra (update status) → Kafka → Chat Server A → sender gets double-tick.

4.6 6. Deep Dives

6.1 WebSocket Connection Management

Why WebSocket instead of HTTP polling?

Method	How	Pros	Cons
Short polling	Client polls every N seconds	Simple	Wasted requests; high latency (up to N sec)
Long polling	Server holds request until message or timeout	Fewer round trips	Still HTTP overhead; server holds threads
WebSocket	Full-duplex TCP connection	True real-time; low overhead	Stateful — servers must track who is connected where

Method	How	Pros	Cons
SSE (Server-Sent Events)	Server pushes, client uses XHR for sends	Simpler than WS	Unidirectional; need separate channel for sends

Key Decision: WebSocket for all real-time delivery.
Connection registry in Redis:

```

1 SET "conn:{user_id}" "chat-server-47" EX 300
  
```

- › Refreshed every 60s by Chat Server via heartbeat.
- › If key expires (server crash), user treated as offline → push notification path.

Load balancing WebSocket connections:

- › L4 load balancer (not L7) — connections are long-lived TCP; no HTTP termination per request.
- › Use consistent hashing on user_id for connection routing (optional but reduces cross-server lookups).

6.2 Message Delivery: Online vs Offline



Message deduplication:

- › Client generates a UUID (client_msg_id) before sending.
- › Server stores client_msg_id → server_msg_id mapping for 24 hours.
- › If retry received with same client_msg_id, return existing server_msg_id (idempotent).

6.3 Group Chat Delivery

Small groups (≤ 500 members):

- › Fanout on message write: write one message record, but create one delivery record per member.
- › Alternative: write one message, each member's client polls for group updates using a last_read_seq cursor.

Recommended for ≤ 500 members:

```

Message write → Kafka
Kafka consumer → write 1 message to messages table
                  → write N "unread" records to user_conversation table (one per member)
  
```


Online members → WebSocket push
 Offline members → Push notification

Very large groups (> 500, e.g., broadcast channels):

- › Switch to fanout-on-read: message is in one place, members' clients pull new messages using seq_num cursor.
- › Similar to news feed pull model.

6.4 Message Ordering and Sequence Numbers

Problem: Two users send messages simultaneously. Who goes first?

Solution: Sequence number per conversation (not global).

- › Each conversation has a monotonic seq_num counter stored in Redis or a coordination service (like Zookeeper).
- › When Message Service processes a message, it atomically increments seq_num for that conversation_id and assigns it to the message.
- › Messages displayed in seq_num order on all clients.

Clock skew: Client timestamps cannot be trusted. Always use server-assigned seq_num for ordering. Client timestamp only shown as display time.

INCR "seq:{conversation_id}" → returns 100034 (atomic in Redis)
 Store message with seq=100034

6.5 Online Presence

Heartbeat protocol:

- › Client sends WebSocket heartbeat every 10 seconds.
- › Chat Server updates Redis: SET "presence:{user_id}" 1 EX 30.
- › Key expires in 30s if no heartbeat → user considered offline.

Displaying presence:

- › When user A opens a chat with user B: fetch GET "presence:{user_id_B}".
- › "Last seen": stored as last_seen:{user_id} on heartbeat stop (WebSocket disconnect or timeout).

Scale challenge:

- › 500M DAU × heartbeat every 10s = 50M presence writes/sec to Redis.
- › **Optimization:** Only send heartbeat if the user is actively viewing a conversation. Idle app = slower heartbeat (30s). Closed app = no heartbeat.
- › Use presence fanout sparingly: only push presence updates to contacts who have a conversation open with that user.

6.6 Read Receipts

States: sent → delivered → read.

sent: Message reached the chat server (ack from server to sender).
 delivered: Message delivered to recipient's device (ack from recipient device).
 read: Recipient opened the conversation (client sends "read" ack).

Implementation:

- › Each message has a status field in Cassandra.
- › Recipient sends {"type":"ack","msg_id":"...", "ack_type":"read"} over WebSocket.

- › Chat Server → Message Service → update Cassandra status → forward read receipt to sender's Chat Server → sender gets green double-tick.

Group read receipts:

- › Store read_by: Map<user_id, timestamp> per message in Cassandra.
- › "Delivered to all" = all members' devices have acknowledged delivered.
- › Show "Read by X, Y, Z" as clients send individual read acks.

4.7 7. Bottlenecks & Scaling

150 million concurrent WebSocket connections:

- › Each Chat Server handles ~50K–100K connections (Go or Node.js event loop).
- › Need ~1,500–3,000 Chat Servers.
- › Chat Servers are stateful but share no business logic — scale horizontally, no coupling.

Kafka throughput (230K messages/sec):

- › Partition by conversation_id (ensures ordering per conversation).
- › 230K msg/sec × 1 KB = 230 MB/s — Kafka handles this easily with dozens of brokers.
- › Message Service consumer group: scale consumers to match Kafka partition count.

Cassandra message storage (3 TB/day):

- › Partition by conversation_id ensures messages in same conversation are co-located.
- › Hot conversation: one partition gets hammered → add bucket (time shard) to partition key.
- › 3 TB/day × 365 = ~1.1 PB/year. Cassandra compresses ~3:1 → ~370 TB/year of disk.
- › TTL old messages: archive messages > 1 year to cold storage (S3 Glacier), keep last 1 year hot.

Media delivery:

- › Don't send media over WebSocket — too large.
- › Client uploads directly to S3 (pre-signed URL from Media Service).
- › Sends media URL in the chat message. Recipients download from CDN-fronted S3.

Push notification scale:

- › 500M users. At any time, 70% offline.
- › Surge: celebrity group sends a message → 1M push notifications simultaneously.
- › Use async workers + APNs/FCM batch APIs. Rate limit per APNs connection (use multiple connections).

4.8 8. Trade-offs Summary

Decision	Choice Made	Alternative	Reason
Real-time transport	WebSocket	Long polling / SSE	Full-duplex; lowest latency; widely supported
Message storage	Cassandra (conv-partitioned)	MySQL (sharded)	Write throughput; time-series access pattern; no joins needed
Connection registry	Redis (TTL-based)	Consistent hash routing	Fast lookup; auto-evicts stale connections on TTL expiry
Message ordering	Server-assigned seq_num (Redis INCR)	Vector clocks	Simple; sufficient for single-conversation ordering
Group delivery	Fanout to members (< 500) / pull (> 500)	Always pull	Balances latency vs write amplification by group size

Decision	Choice Made	Alternative	Reason
Presence heartbeat	10s client heartbeat, 30s TTL	Server-push disconnect events	Handles silent disconnects (network failure, app killed)
Read receipts	Per-message status in Cassandra	Separate receipt table	Avoids join; co-located with message data
Media	S3 + CDN (URL in message)	In-message blob	Media over WebSocket is impractical at scale

End of Part 1 — Foundational System Design Problems.